

Courting the Academy

The Judicial Role in Popularizing Legal Scholarship*

Matthew Estes[†] Ransi Clark[‡]

August 1, 2025

Abstract

The relationship between the courts and legal academia is a two-way street. Legal scholars propose legal theories that judges use to support their opinions. Soon other academics join the fray, giving their own assessments of the judicial use of the theory, raising the academic profile of the legal theory and the legal scholar. Selective in their acknowledgments, judges become not just arbiters of justice but also academic success. The judge-academic dynamic, though widely noted, is under-studied empirically. This paper closes the gap by introducing a causal framework to explore how legal theories become legal canon.

Combining data on academic citations and judicial mentions, we quantify the extent to which judges impacted the popularity of two legal movements: legal realism and originalism. Exploiting timing differences in judicial mentions between legal works, we apply a conservative bounding approach to estimate the degree to which academic popularity is spurred on by judicial acknowledgement. We find that judicial mentions increased academic citations to legal realist works but not for originalist works. Although the two movements are not contemporary and therefore the findings cannot be used to compare the relative judicial role in the movements net of other factors, the empirical exercise affirms the judicial-academic relationship.

Keywords: empirical legal studies, courts, realism, originalism, counterfactuals

*This research uses data made available through JSTOR's Constellate platform, operated by ITHAKA. The findings and conclusions are solely those of the authors. We are grateful to Mike Alvarez, Danny Ebanks, Jonathan N. Katz, Ryan T. Moore, Jacob Morrier, Ronak Patel, Bob Sherman, and participants at the 2025 MPSA Conference for helpful comments and feedback on this project.

[†]A.B. 2018, Harvard University; J.D. 2021, The University of Chicago Law School; Ph.D. Candidate, California Institute of Technology. mestes@caltech.edu

[‡]B.A. 2018, University of Queensland; Ph.D. Candidate, California Institute of Technology. ransi@caltech.edu

1 Introduction

The Supreme Court’s recent decision in *Franklin v. New York*¹ offers a glimpse into the contemporary relationship between the judiciary and the legal academy. In *Franklin*, the Court denied *certiorari* to reconsider Confrontation Clause jurisprudence under the Sixth Amendment. Yet Justices Alito and Gorsuch issued opinions signaling a willingness to overturn a leading Confrontation Clause case.² Why? Because, Justice Alito writes, “recent scholarship ... casts doubt on key aspects of *Crawford*’s reasoning.”³

The Court’s encouragement of particular legal scholarship is not new. Indeed, a large amount of empirical legal scholarship studies how judges use legal scholarship. What is less studied is the academic reaction to judicial acknowledgment. This reverse relationship is just as important: despite criticisms that the legal academy is disengaged with the practical matters of lawyering and the judiciary, individual academics may hold judicial mentions in high esteem. Academics with judicial ambitions can be incentivized to produce work that they believe the judiciary will invoke. To provide for the judiciary, academics attend to the judiciary’s preferences. Judges can signal a need for further academic work by mentioning certain legal pieces in their opinions. Other judges may produce the scholarship themselves.⁴ Still others may convey signals by indirect means, e.g. through their clerks who later become academics.

The purpose of this paper is quantification of the extent to which the academy heeds the judiciary’s call. More specifically, we consider the courts’ role in the rise of legal realist scholarship in the early 20th century and the rise of originalist scholarship in the late 20th century. We chose these specific academic movements for their historical and contemporary relevance and also because their adherents and detractors are (relatively) readily identifiable.⁵ The objects (or units) of analysis are scholarly works (books, chapters, or articles) that are most relevant to legal realism or originalism. For each work, we collected time series data on judicial mentions⁶ and an annual series of academic citations. Those specific series are obtained with a view to esti-

¹*Franklin v. New York*, 604 U.S. (2025).

²*Crawford v. Washington*, 541 U.S. 36 (2004).

³*Franklin v. New York*, 604 U.S. (2025) (Alito, J., concurring).

⁴Consider Judges Posner and Easterbrook, or Justice Scalia.

⁵Although there are a great many attempts by academics to launch legal movements, it is difficult to round up all such attempts and/or ascertain the boundaries of particular movements.

⁶We use the term “judicial mentions” instead of the more familiar “judicial citations” to avoid confusion between citations arising from the judiciary and the academy.

measuring how much of scholarly attention is attributable to court mentions. Although citations are an imperfect measure of scholarly attention, they are an oft-used proxy for academic popularity.

In contrast to citation analyses which explore the divergence between the academy and judiciary,⁷ our aim is to investigate the more specific question of whether academic citations respond to judicial acknowledgment. The reason to choose such a specific aim is so that we may employ the machinery of causal inference and make causal claims, albeit with caution. In general, a causal relationship between judicial mentions and academic citations is difficult to pin down. For example, the two series may be simultaneously affected by the politics of judicial appointment.⁸ Academics who see the political tides changing may get ahead of it by producing scholarly work that is adherent to a particular judicial movement.⁹ Despite these difficulties, there is much variation in how¹⁰ (and how often) works are mentioned by the judiciary, with some works mentioned much more often than others.

Challenges to causal inference are encountered across the social sciences. Because of the insolubility of many such challenges, there is a preference towards designing and implementing experiments whenever possible. Unfortunately, the question at hand is not easily amenable to experimentation. Such is the case for many questions concerning the judiciary, not least because diversions from normal operations even for scientific purposes may be seen as an illegal or unethical miscarriage of justice.¹¹ Yet, experimentation difficulties should not be fatal to the study of these processes, provided the counterfactual assumptions underpinning the interpretation of the derived estimates as causal hold. At its core, the counterfactual method here separates works into two groups:¹² a treatment group (works that mentioned by the judiciary) and a

⁷See [Merritt and Putnam \(1995\)](#) for an example of such an analysis.

⁸Both realist and originalist movements were colored by the political leanings of the Administrations that nominated adherent judges to the court. The Roosevelt Administration nominated several legal realists and the Reagan Administration nominated originalists.

⁹They may also do so in anticipation of a judicial appointment or a blockbuster court case.

¹⁰Judges and their clerks may also use, for instance, citation counts as a heuristic in determining whether the work should be mentioned in their opinion or not. This potentially summons the pernicious problem of reverse causality. The possibility is not that the judicial mentions increased the scholarly popularity of the work, but its own baseline level of scholarly popularity.

¹¹Cf. [Ho \(2005\)](#) at 1999 (“The problem for legal scholars and social scientists is that laboratory experiments are often infeasible, expensive, or unethical.”)

¹²Separation of units into treatment and control groups might not be straightforward, particularly in dynamic settings such as ours. The process requires coarsening of the paths since some works receive a disproportionately large number of mentions. There may also be too few works that receive no mentions from the judiciary, even though there will be a large number of works that receive a very small number of mentions dispersed over several decades.

control group (works not mentioned).¹³ The specification of such groups is necessary since it is the citation outcome of the control group that is used to construct counterfactual citations for a treatment group under alternative (i.e. counterfactual) judicial mention scenarios. But what counterfactual assumptions should one use here?

Since different causal assumptions lead to different estimates, we choose a bounding method that nests several of these causal assumptions. Our method,¹⁴ which we call the bounded deviations approach, allows for a wide range of counterfactual outcomes in citation trajectories. Despite the statistical jargon, the counterfactual intuition is simple: absent treatment, treated units would’ve experienced an outcome “close to” or “nearby” the control outcomes. The upshot is that, rather than reporting a result like “the causal effect is X ” under strong counterfactual assumptions, we report a range (“the causal effect is between X and Y ”) under weaker assumptions.¹⁵ An important advantage of this bounding method is that it renders itself to sensitivity analysis almost immediately: do the conclusions change if confidences in the counterfactual assumptions change?¹⁶

¹³Alternatively, in some comparisons we can compare works mentioned early versus those mentioned late, if the timing is plausibly exogenous. See *infra* [Section 4](#) for further discussion.

¹⁴Our approach follows the partial identification literature in econometrics, especially that of Manski. See, e.g., [Manski \(1997\)](#), [Manski \(2007\)](#), [Manski and Pepper \(2018\)](#).

¹⁵An advantage of bounding is in applications such as ours where there is uncertainty as to whether the standard assumptions required to make causal claims are satisfied. For example, the parallel trends assumption in differences-in-differences analysis, which allows the causal effect to be the post-treatment difference less the pre-treatment difference immediately before treatment. Another source of assumption uncertainty in our setting arises from the dynamic nature of the data: for example, the uncertainty arising from the process that generates academic citations to texts overtime, or uncertainty arising from unknown treatment dynamics.

Another consideration is the sequential nature of treatments in our study. The judicial mentions treatment studied here is an entire sequence of events. The impact of this sequence or treatment history is not necessarily immediate within the year. A treatment that is a sequence of events rather than one event can cascade over time too. A bounding approach is more appropriate here, because it makes better use of the pre-treatment behavior of the time series.

The approach is flexible compared to the regression implementation for the textbook differences-in-differences. Yet, limitations exist even within this more flexible framework because the more varied the treatment paths are the more numerous their possible counterfactuals become. Statisticians call this problem the curse of dimensionality. See discussion below.

¹⁶Since the bounds are derived from multiples of the pre-treatment difference, changing the multiple can determine how many multiples of the pre-treatment difference is required before the treatment effect’s direction becomes inconclusive. Statistically minded readers will wonder if confidence intervals on point estimates perform the same function as the bounds we suggest here. The answer is no, because confidence intervals are produced by assuming that the observed data is a sample from a larger population and is chosen to reflect this sampling error. We do not suppose that the chosen texts come from a larger sample, even though our choices are not exhaustive. The main source of uncertainty here is that of what assumption best reflects the pre-treatment behavior of the two series.

See generally [Manski \(1993\)](#) at 4:

Using the bounded deviations approach, we find that judicial mentions increase academic citations to realist works. These effects generally appear with some delay and are relatively robust to the identification tuning parameter we use. We show, for example, that realist works first mentioned in judicial opinions from 1976–2025 received a positive boost in academic citations. Although the magnitude of the effect is not identified, it is between 5–25 added academic citations across model specifications.

The impact of judicial mentions on originalism is less certain. Because it is a nascent judicial theory, it is too early to tell whether academic citations to originalist works increase or decrease in response to judicial mentions. The lag between issued judicial opinions and the development or notice of academic legal theories suggest the data will not pin down the effect for some time. Evidence of causality (or lack thereof) may, in other words, be premature because an uncertain amount of time may need to elapse for an academic work to trickle into the judicial and academic milieu.¹⁷

Our work contributes to several strands of literature. The most established of these is that of the origin of jurisprudential movements. We bring citation metrics to bear on the evolution of such movements. The beginnings of jurisprudential movements, in particular those that have partisan bearings, are usually prone to revision¹⁸ and reinvention.¹⁹ The process of identifying influential texts and then tracing their influence on the academy and the judiciary is less prone to such temptations. Although legal movements are often analyzed from an historical viewpoint, we hope to bring quantitative methods to legal origin stories by considering the judiciary’s role in the rise and fall of legal movements.

Another strand of literature, beginning with Judge Harry Edwards’ observation that “law and” movements are displacing practical legal training in law schools,²⁰ seeks to quantify the relationship between the legal academy and the judiciary with

“It is useful to separate the inferential problem into statistical and identification components. Studies of identification seek to characterize the conclusions that could be drawn if one could use the sampling process to obtain an unlimited number of observations. Studies of statistical inference seek to characterize the generally weaker conclusions that can be drawn from a finite number of observations. Statistical and identification problems limit in distinct ways the conclusions that may be drawn in empirical research. Statistical problems may be severe in small samples but diminish in importance as the sampling process generates more observations. Identification problems cannot be solved by gathering more of the same kind of data.”

¹⁷Because little of the mechanism by which judges learn of legal scholarship is known (except for their own pronouncements), it is difficult to know when it is too early to assess the causal impact of judicial (non-)mention on the fate of the publication.

¹⁸See [Tamanaha \(2008\)](#) for realism and [Sawyer III \(2018\)](#) for originalism.

¹⁹Both legal realism and originalism has “new”-er versions.

²⁰[Edwards \(1992\)](#)

a view to studying the divergence between the two institutions. The vast majority of such empirical studies quantify the overall volume of judicial mentions of legal scholarship. While volume is an important metric and adequate for an exploration of whether the judiciary uses legal scholarly work, it does not give a complete picture of the relationship: by focusing on particular substantive movements, we are better able to measure the effect of judicial mentions on subsequent scholarly work relevant to that movement.

Lastly, we use this empirical setting to demonstrate a transparent causal method in dynamic settings. Bounding is robust to a wide range of assumptions researchers make to pin-point a specific causal quantity, many of which go unstated. By proposing a wider range, we reduce the need for numerous or implausible assumptions. The assumption that we make—that the post-treatment difference between two counterfactual series is no more than some multiple of the pre-treatment difference—is concise, transparent, and nests other popular approaches.²¹

The remainder of the paper is structured as follows. In [Section 2](#), we review existing empirical and non-empirical work on how courts use academic legal theories and works. We also discuss how existing scholarship largely neglects how the courts affect legal academics. In [Section 3](#), we give an overview of two important U.S. legal movements: American Legal Realism and Originalism. We discuss the observational data used in our study and some high-level descriptive findings from the data. In [Section 4](#), we discuss the causal inference problem in this setting. We describe simplifying assumptions and illustrate how our causal strategy works. Finally, in [Section 5](#), we apply our methodology to estimate the impact of judicial mentions on academic citations to realist and originalist works.

2 Legal Theories in the Academy and in the Courts

The relationship between the legal academy and the courts is a subject of interest for many legal scholars. Still, there is much descriptive empirical work that remains undone. Although there has been much hand-wringing about how academics affect (or not) judges and the judicial process, the reverse channel is relatively underexplored. In

²¹For example, our approach nests the parallel trends assumption for the difference-in-differences (“DID”) strategy. See, e.g., [Manski and Pepper \(2018\)](#). In the DID approach, comparisons are made between citation trajectories re-centered at their baselines. This can inoculate against a variety of challenges to causal claim-making such as: popularity of the academic work (or the author) before judicial mentions, citation trends that are affected by factors extraneous to judicial mentions, or citation trends related to the time elapsed since publication.

particular, credible empirical evidence of the causal effects of judicial activities on the development, production, and reception of academic legal theories is not available. This section discusses prior work—both empirical and non-empirical (qualitative, anecdotal)—on how the courts and legal academy interact with one another.

2.1 How Courts Use Academic Legal Articles

One well-known lament is that legal academics are largely irrelevant to judges and judging. There is a smattering of qualitative and anecdotal evidence for this claim. Consider, for example, Judge Posner’s claims along these lines:

“So there is much to criticize in the judicial profession and therefore much room for improvement. But where is the improvement to come from? A possibility that appeals to me as a former law professor is the law schools. Law professors write a great deal about the judiciary—and mainly the federal judiciary. But there is a question about how well informed about, or helpful to, the judiciary that writing is. At present, not very, I have discovered. Not that I’m inclined to apply Brendan Behan’s comment about film critics to law professors by comparing the professors (relative to judges) to ‘eunuchs in a harem; they know how it’s done, they’ve seen it done every day, but they are unable to do it themselves.’ Some judges might think the comparison apt, however.”²²

Empirical findings support this hypothesis. At the Supreme Court, for example, [Newton \(2012\)](#) finds that only 20.19% of signed opinions mention a law review article from 2001-2011.²³ This may, in part, stem from ideological differences between the justices and the authors of law review pieces: the three justices with the highest percentage of opinions mentioning law reviews were liberal justices (Breyer 26.14%, Stevens 24.54%, and Ginsburg 23.08%), whereas the three justices with the fewest citations to law reviews were all conservative justices (Thomas 13.31%, Rehnquist 13.79%, and Roberts 16.04%).²⁴

This trend may be changing. Although the number of mentions to law review pieces is generally low,²⁵ there is some evidence that law review pieces are increasingly

²²[Posner \(2016\)](#).

²³See Table 1 in [Newton \(2012\)](#).

²⁴See Table 1 in [Newton \(2012\)](#).

²⁵[Sirico Jr. and Drew \(1991\)](#)

mentioned, especially elite law reviews in high-profile Supreme Court cases.²⁶ Indeed, [Detweiler \(2020\)](#) finds that, although mentions of law review articles remained low in 2018 (1.8%), opinions increasingly mention law reviews relative to the historical minimum in 2009 (1.47%).²⁷ Note that the [Detweiler \(2020\)](#) study uses the entire universe of reported opinions from U.S. state and federal courts 1945–2018.²⁸

Another important factor affecting judicial mentions of academic (and non-academic) sources is technological progress. [Fronk \(2010\)](#) analyzes federal appellate citation patterns and finds that the average number of cited cases increased from 15.66 in 1957 to 31.14 in 2007.²⁹ This was a period of technological progress, which substantially lowered the cost and difficulty of locating relevant legal sources. A similar effect may be relevant for citations to law review works, which are increasingly easy-to-access and widely shared, if law review pieces are *complements* to primary legal sources.³⁰

A final consideration is that legal opinions may cite law review articles because they speak to a judge’s own theories. Legal theories may not constrain judicial decisions writ-large, but they may persuade an individual judge to the extent they align with the judge’s own theories.³¹ Indeed, judicial mentions of more jurisprudential or theoretical material may represent alignment with or acceptance of theoretical precepts, means, and ends. To the extent law review works better reflect the theories of practicing judges, we may expect more citations to corroborating law review works.³²

²⁶See [Feldman \(2018\)](#).

²⁷[Detweiler \(2020\)](#) attributes the overall decline from the peak in early 1980s to substitution with primary legal sources.

²⁸What is driving this possible shift remains unclear. One possibility is that law review articles are becoming more sophisticated and empirical, which can “bolster a doctrinal claim.” See [Diamond and Mueller \(2010\)](#) at 595. Indeed, the amount of “minimal empirical content” in law review articles increased from 1998–2008. See Figure 3 in [Diamond and Mueller \(2010\)](#).

²⁹See [Fronk \(2010\)](#) at Table 1.

³⁰But see [Detweiler \(2020\)](#) for discussion of law review works and primary legal sources being *substitutes*. In this case, the technological progress for easily accessing primary legal materials would decrease citations to law reviews. Indeed, [Detweiler \(2020\)](#) argues that: “By facilitating the discovery of primary law, these services provided researchers with a viable, and in many ways superior alternative to using law reviews as a shortcut for their own research.”

³¹See, e.g., [D’Amato \(1999\)](#); see also [D’Amato \(1989\)](#).

³²If this explanation is correct, should we expect further increases in citations to law reviews as machine learning and generative language models perform increasingly well at legal reasoning tasks and law professors (and lawyers) are better able to pin down a judge’s operant theories? See, e.g., [Thalken et al. \(2023\)](#) (discussing LLMs and legal reasoning tasks).

2.2 How Courts Affect Legal Academia

Even if judges find law reviews or legal academia “irrelevant”, this does not imply that judges are irrelevant to law professors. Indeed, judges may impact the development in many ways, intentional or otherwise.

The most obvious influence of the courts on legal academia is on legal educational materials. Because of the casebook method, judges impact legal education through written opinions. Empirical descriptive work suggests that certain superstar judges (e.g. Judges Richard Posner and Frank Easterbrook) dominate the casebooks.³³ Longer service (i.e. seniority) is also positively correlated with casebook inclusion.³⁴ Other studies suggest that additional factors may explain casebook inclusion, especially for non-senior judges with a high degree of casebook inclusion.³⁵

Another way judges can impact law professors is by changing the demand for certain legal theories. Simple economic incentive arguments imply that academics respond to the non-pecuniary benefits—e.g. prestige, fame, citations, interest—from mentions in court opinions. Indeed, increased court demand (support) for particular legal theories should, on balance, increase the supply of those theories unless law professors and judges operate in fundamentally different “markets.”. Given the influx of scholars from other disciplines into the legal academy, academics may feel more pressure to impress like-minded colleagues rather than communicate to judges on the practical matters of the law. For some judges and some law professors, this is certainly a possibility, as Judge Posner notes:

“[T]here is a wild literature that I have avoided mentioning in which law professors in immensely long articles subject legal texts to the hermeneutic techniques of postmodernist literary theory. No judge could get *anything* out of that literature, and this unbridgeable gap is not merely a generational one.”³⁶

Judge Posner offers an especially vivid potential explanation for this phenomenon:

“The process is Darwinian. In nature each animal species must find a niche for itself, critically including a food that it can find and eat without encountering destructive competition from another species. In the

³³Gulati and Sanchez (2002).

³⁴Gulati and Sanchez (2002).

³⁵See, e.g., Fitzpatrick and Varghese (2017).

³⁶Posner (1992) at 1928.

academy each species of professor must find an academic niche in which he can avoid destructive competition from other professors. ... Their need to communicate with persons outside of their niche, such as judges, like the need of a squirrel to learn to eat dandelions as well as nuts, is minimized.”³⁷

In other words, there is nothing inherently “wrong” with market specialization: judges and law professors simply occupy different niches.

Politics and ideology are also likely relevant to whether legal scholars respond to courts. Because legal scholarship is often ideological,³⁸ conservative legal scholars may write, for example, to respond to the current Supreme Court whereas liberal legal scholars will not. In other words, scholars may operate in separate ideological silos, which can affect whether certain law professors respond (or not) to court demand.

The focus of the remainder of this paper is testing whether law professors do, in fact, respond to judicial mentions. We examine whether certain academic law pieces increase in popularity following mentions in published court opinions. Because popularity is measured in academic citations to the original piece, we are measuring whether judicial mentions have a particular role in spurring on academic movements. To avoid both problems of over- and under-inclusion,³⁹ we focus our attention on two specific jurisprudential movements and obtain as many relevant pieces that are recognized as within the movement, including those of their critics. We turn now to describing the two movements and the data we collected.

3 Data & Empirical Descriptions of the Movements

This paper uses data on two important legal movements: American Legal Realism and Originalism. We study these two legal movements because of their historical and contemporary impact on the law. As Posner (2016) notes, realism and originalism are two of the most influential legal philosophies in the United States:⁴⁰

³⁷Posner (2016) at 8-9.

³⁸See, e.g., Table 3 and Figure 3 in Chilton and Posner (2015).

³⁹The ideal population here would be the universe of all legal pieces that forward a jurisprudential theory and are therefore relevant to judging. Given this set, one could make comparisons among the pieces mentioned by the judiciary and those that were not. But such a universe of works is difficult to obtain because there are a great many of them. Another option is to randomly sample scholarly works and determine their relevance manually. But this might be even more unsatisfactory than the indiscriminate use of the entire sample if the random sample omits highly relevant or influential scholars.

⁴⁰As Posner (2016) notes, the origins of these philosophies were also political.

“Legal realism...was a highly influential judicial philosophy not just in the 1920s and 1930s but in the entire period that began with the publication of Holmes’s *The Common Law* in 1881 and ended with the end of the Carter Presidency a century later. Reagan’s election, his conservative appointments to the federal courts of appeals and the Supreme Court, and the more or less simultaneous creation of the Federalist Society, began to shift the federal judicial balance back toward formalism, giving rise to ‘originalism’ and ‘textualism’ and increased resort to dictionaries (‘literalism’) as sources of statutory meaning—all backward-looking sources of judicial guidance.”⁴¹

Because of the ideological character of these movements, their endorsement was likely to bring greater rewards.⁴² Judges are likely to signal their endorsement by mentioning prominent scholarship of the movement. The mention of these by the courts can have an encouraging effect on legal academics who seek to join the fray.

To assess the impact of courts on these two movements, we collected two main types of data for each movement. The first type of data is data on judicial mentions, that is, how often judges cite specific legal works (articles) for thinkers associated with the movement. The second type of data is information on the popularity of different legal works, as measured by academic citations. The data sources, collection process, and high-level descriptive findings are summarized for each movement below.

3.1 Legal Realists

3.1.1 Description of the Movement

American Legal Realism began coalescing as a movement in the 1920s and came to prominence during the Roosevelt Administration, when it began to greatly impact the development of the law in the United States. For legal education and scholarship, the realists’ primary contribution was in developing a non-formalist theory of adjudication. The formalism or “mechanical jurisprudence” of the day held,⁴³ roughly, that judges decide cases by legal rules and reasons.⁴⁴ By contrast, the realists held

⁴¹Posner (2016) at 86-87.

⁴²In the case of originalism, for example, this could entail promotions for individual judges endorsing originalism. Note, however, that Judge Posner is a prominent exception, given his criticism of originalism despite being a Reagan appointee.

⁴³See Pound (1908).

⁴⁴Leiter (2005).

that legal rules and reasons were (in at least some cases) window-dressing for judicial decision-making. Judges actually decide cases based on the facts of the case. As [Leiter \(2010\)](#) puts it: “judges are largely ‘fact-responsive’ rather than ‘rule-responsive’ in reaching decisions.”⁴⁵ Realists suggested, moreover, that the law was, at least in “hard” cases, rationally indeterminate: the applicable legal rules or reasons did not uniquely determine a particular outcome.⁴⁶

Although all realists accepted these central tenets, they differed substantially on what determines how judges respond to facts. [Leiter \(2010\)](#) describes the split as two wings: the Idiosyncrasy Wing and the Sociological Wing.⁴⁷ Broadly, the Idiosyncrasy Wing (typified by Jerome Frank) held that the individual personality, ideology, psychology, and other background characteristics of judges determine how they respond to facts. By contrast, the Sociological Wing held that certain general social facts determine judicial outcomes. The difference is largely, therefore, one of degree.⁴⁸

As [Leiter \(2005\)](#) observes, the realists were proponents (in theory, if not in practice) of scientific methods of discovery and empirical testing.⁴⁹ Because most realists thought certain general social facts determined judicial outcomes, many were interested in discovering which factual patterns result in certain judicial responses. Some of the realist works then are, unsurprisingly, descriptive empirical projects to specify how particular fact patterns result in certain types of judicial decisions.

Although largely a descriptive theory of adjudication,⁵⁰ the realists also impacted the development of “law in action” and the practice of lawyering. For example, realists worked to reform, improve, and summarize the law of contracts. Corbin, for example, wrote his famous treatise on contract law, whereas Karl Llewellyn was instrumental in drafting and promoting the Uniform Commercial Code. Other realists served as judges (Cardozo, Frank, Holmes, Posner) and in executive agencies (Frank, Cohen).

3.1.2 Data Collection

Although many scholars undoubtedly contributed to legal realism, we focus here on the more foundational and influential realist works. Specifically, we examine 16

⁴⁵[Leiter \(2010\)](#) at 249; cf. at 257.

⁴⁶See also [Leiter \(2005\)](#) (discussing how realists also subscribed to a *causal indeterminism* thesis).

⁴⁷See [Leiter \(2010\)](#) at 257-262.

⁴⁸Social scientists today might say the split was in terms of the degree of heterogeneity in the social facts that affect judicial decision-making.

⁴⁹See [Leiter \(2005\)](#) at 50-51.

⁵⁰See [Leiter \(2010\)](#).

prominent legal realists and their major works.⁵¹ The 16 writers and their major works are included in [Table 1](#).⁵² For each work, we then created two annual time series: (1) academic citations, and (2) judicial mentions.

⁵¹These works were chosen by using the primary sources in survey pieces on American Legal Realism. See, e.g., [Leiter \(2005\)](#) and [Leiter \(2010\)](#). Also, note that although Holmes' *The Common Law*, which was mentioned above in the Posner quote was surely influential, the name of the work is too common to obtain a reliable academic citation series.

⁵²We note that many of these works are often considered to be canon. See, e.g., [Kennedy and Fisher \(2006\)](#).

Table 1. Legal Realists and Their Works

Writer	Work
Arthur Corbin	Corbin on Contracts
Benjamin Cardozo	The Nature of the Judicial Process
Benjamin Cardozo	The Paradoxes of Legal Science
Felix Cohen	Transcendental Nonsense and the Functional Approach
Herman Oliphant	A Return to Stare Decisis
Jerome Frank	Are Judges Human?
Jerome Frank	Law and the Modern Mind
John Chipman Gray	The Nature and Sources of the Law
Karl Llewellyn	A Realistic Jurisprudence
Karl Llewellyn	Remarks on the Theory of Appellate Decision and the Rules or Canons about How Statutes are to be Construed
Karl Llewellyn	Some Realism about Realism
Karl Llewellyn	The Bramble Bush
Karl Llewellyn	The Common Law Tradition: Deciding Appeals
Leon Green	The Duty Problem in Negligence Cases
Leon Green	The Judicial Process in Torts Cases
Max Radin	In Defense of an Unsystematic Science of Law
Max Radin	Law as Logic and Experience
Max Radin	Statutory Interpretation
Max Radin	The Theory of Judicial Decision: Or How Judges Think
Morris Cohen	Property and Sovereignty
Morris Cohen	The Basis of Contract
Oliver Wendell Holmes	The Path of the Law
Richard Posner	The Problematics of Moral and Legal Theory
Robert Hale	Force and the State
Roscoe Pound	Mechanical Jurisprudence
Roscoe Pound	The Scope and Purpose of Sociological Jurisprudence
Underhill Moore	An Institutional Approach to the Law of Commercial Banking
Underhill Moore	Law and Learning Theory: A Study in Legal Control
Underhill Moore	Legal and Institutional Methods Applied to the Debiting of Direct Discounts
Underhill Moore	Rational Basis of Legal Institutions
Wesley Hohfeld	Some Fundamental Legal Conceptions as Applied in Judicial Reasoning

To produce the academic citations series, we use data from the JSTOR Data for Research (DfR) program. This data contains information on the articles that cite the author’s last name and the title of the work. The outcome citations are from academic articles only and not, e.g., popular coverage of the work. Because some of the works have very general titles (e.g. “Statutory Interpretation”), the search we carried out is for the author’s last name and an exact match to the title of the realist work to avoid false positives.

The second time series is how many times the courts mention realist works, which serves as our “treatment” path data.⁵³ To collect judicial mentions, we use public data from CourtListener.⁵⁴ This data is then used to encode treatment paths for each work, which represent the number of times a particular work is mentioned each year. For example, the treatment path $\vec{k} = (0, 0, 1, 1, 2)$ means a work receives zero court mentions in years 1 and 2, one court mention in years 3 and 4, and two court mentions in the final year. To avoid false positives in judicial mentions, the search we carried out was for the author’s last name and an exact match to the title of the realist work.

The cumulative counts aggregated across realist authors is shown in Figure 1. The left panel is the cumulative JSTOR citations series over time. This is the cumulative number of works on JSTOR that cite to one of the Table 1 realist works. The right panel is the number of court mentions. This is the cumulative number of published judicial opinions that mention one of the realist works.⁵⁵

⁵³See also below Section 4.2.

⁵⁴See <https://www.courtlistener.com/>

⁵⁵In the appendices, we also show plots for individual authors. In Figure A1, for example, we plot the cumulative citation series for each work separately. In each author panel, the citation series are shown for the works in Table 1. We also show the citations on a yearly, non-cumulative basis in Figure A2. In Figure A3, we show the yearly court mentions for each author (panels) and work (lines). For example, in 1996 Arthur Corbin’s *Corbin on Contracts* was mentioned in 79 court cases. By contrast, Holmes’ “The Path of the Law” and Cardozo’s “The Nature of the Judicial Process” were each mentioned 8 times by courts in 1996.

Cumulative Realist Counts

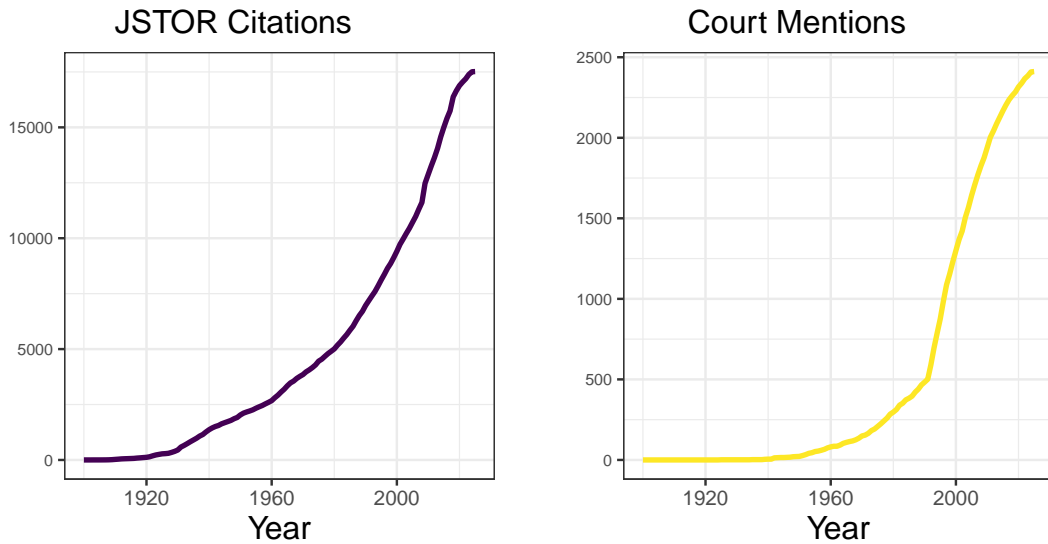


Figure 1. Cumulative Realist Counts

Note: The y -axis differs by panel. Academic citations are more numerous than court mentions for every year.

The works receive court mentions at different times and levels. In [Figure 2](#), we show that some works are highly mentioned early on relative to their publication date (e.g. *Corbin on Contracts*). By contrast, some works only receive judicial notice many years after initial publication (e.g. Pound’s “Mechanical Jurisprudence”) or are never mentioned by the available CourtListener opinions (e.g. the works of Underhill Moore).⁵⁶ That so few works share similar treatment paths (i.e. treated at the same time and level) is important to the method we use to identify causal effects. Because our method can be used on individual works,⁵⁷ the fact that there is little overlap in treatment paths requires modifications to simple causal inference strategies.

⁵⁶See [Appendix C](#) for full court mention histories.

⁵⁷As in, e.g., the synthetic control literature pioneered by [Abadie et al. \(2010\)](#).

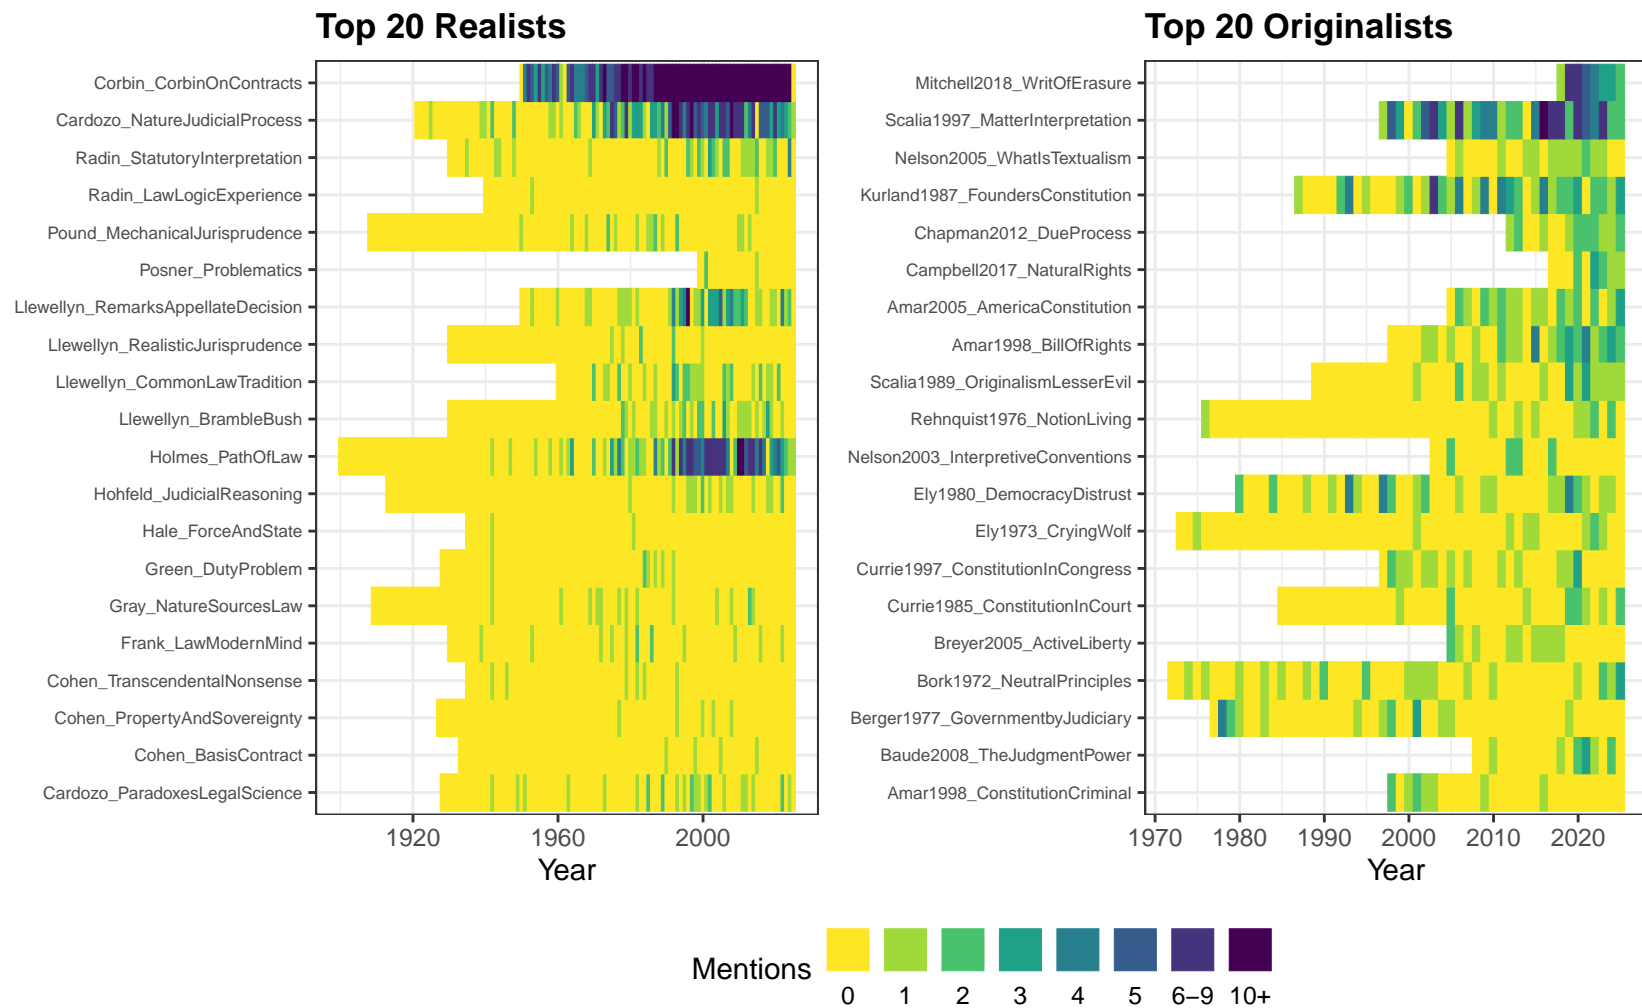


Figure 2. Realists & Originalists: Court Mentions

3.2 Originalists

3.2.1 Description of the Movement

Originalism solidified in the 1980s as a movement and is one of the more influential (and controversial) legal movements of the modern era. The most basic tenet of originalism is that the authoritative meaning of constitutional text is the meaning at the time of enactment.⁵⁸ This proposition, originalism proponents argue, restrains judges, restricts the permissible readings of ambiguous constitutional texts, and is “our law.”⁵⁹

Originalism is unlike legal realism in at least two important ways. First, it is concerned exclusively with constitutional law. The realists, by contrast, studied across legal subjects and were particularly interested (as academic and practitioners) in commercial law.⁶⁰ Second, whereas realism was largely a descriptive theory of adjudication, originalism advances a normative and prescriptive theory for how judges should decide constitutional law issues.⁶¹ Accordingly, many originalist works advance particular “originalist readings” of a constitutional clause and argue that is how judges should decide cases involving that clause.

Although originalism began as a contested descriptive and normative academic theory, its influence on the judiciary is now widely appreciated. Originalist theory has impacted how judges justify their decisions in constitutional cases.⁶² And the rise in judicial acceptance of originalism has generated more demand for originalist scholarship. This is an almost uniquely American phenomenon: originalism in constitutional theory is not widely embraced outside the United States.⁶³

3.2.2 Data Collection

Unlike legal realism, which was largely associated with two law schools (Columbia and Yale),⁶⁴ originalism claims many varied adherents across law schools with methodological and substantive differences. Because it is also an active and ongoing legal

⁵⁸See, e.g., [Whittington \(2004\)](#).

⁵⁹[Baude \(2015\)](#)

⁶⁰See above for some discussion of realism and contract law.

⁶¹But note that some originalists also argue for the descriptive claim that originalism is the law. See [Baude \(2015\)](#).

⁶²For example in *District of Columbia v. Heller*, both the majority and dissenting opinions devote countless pages to uncovering the 18th century meaning of the Second Amendment.

⁶³See [Greene \(2009a\)](#) and [Greene \(2009b\)](#).

⁶⁴[Leiter \(2005\)](#) at 51.

movement, the foundational originalist authors and texts are not as clear as in legal realism. Accordingly, we have included many more possible foundational originalist works than realist works.

To construct our population of originalist texts, we used a syllabus for an originalist class (“Originalism and Its Discontents”) at the Harvard Law School taught by Professor Stephen Sachs.⁶⁵ The syllabus contains a long list of optional readings relevant to originalism—and, therefore, allowed us to cast a wide net in our search for “foundational” originalist texts. We omitted all non-academic articles (e.g. court cases) that are listed on the syllabus, as well as works for which we could not find academic citation series in JSTOR, or for which the titles of the work were so common (e.g. “Commerce Clause”) that they would not reflect citations to the academic article alone.

Our final list contains 191 originalism-relevant texts. As in the realist case, the list of works is neither exhaustive nor definitive. Nevertheless, it includes many popular and influential originalist texts. Importantly, note that not all of the works are by authors that adhere to originalism. Some of the works are merely *relevant* to originalism because this is a course syllabus.

As before, we generated two time series for each originalist work. The first is the number of academic citations to each work (the JSTOR citation series). The second series is how many opinions each originalist work is mentioned in each year. The cumulative time series are plotted in Figure 3: the cumulative JSTOR citations series aggregated across authors is displayed in the left panel; the cumulative court mentions series aggregated across authors is displayed in the right panel.

⁶⁵We used the Spring 2025 Sachs syllabus available online at <https://www.stevesachs.com/syllabi/originalism.pdf>.

Cumulative Originalist Counts

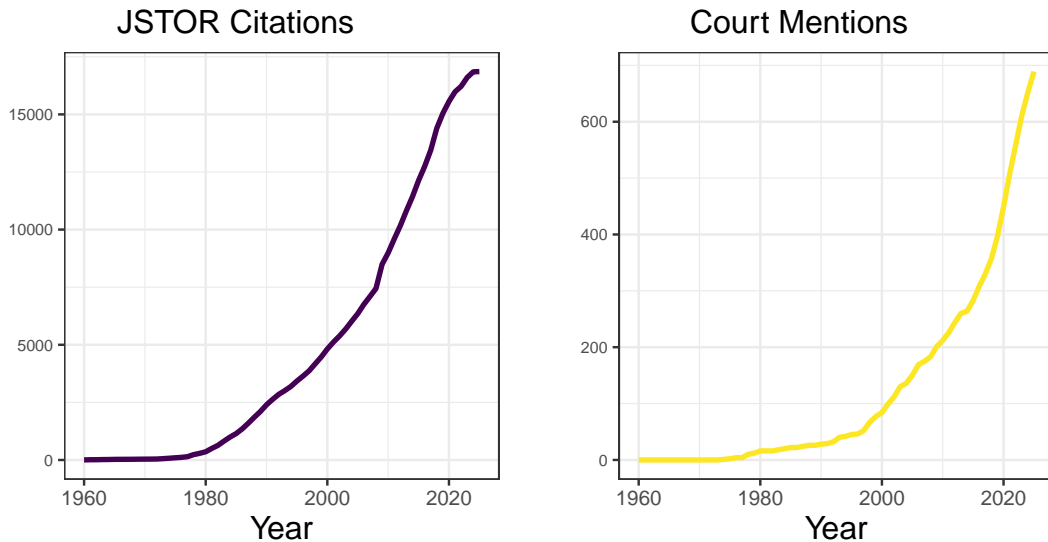


Figure 3. Cumulative Originalist Counts

Note: The y -axis differs by panel.

As in the realism case, the originalist works are mentioned by judges to varying degrees and at varying times. The right panel of [Figure 2](#) also shows the levels and years the different works are mentioned in published judicial opinions for top originalist works.⁶⁶ As above, it is important to notice here that the “treated” works experience little overlap in the treatment path space.⁶⁷ Unlike the realist case, however, there are many works that are never treated (i.e. never mentioned in a judicial opinion). This is not unexpected because originalism is a newer legal theory than realism.

3.3 Descriptive Findings

3.3.1 Textual Data & LLM Classification

This paper is primarily concerned with the academic citation and judicial mention metrics. However, we note here briefly that the JSTOR data in particular contains additional textual data as well that may be of interest for text analysis, machine

⁶⁶The timelines for all works are shown in [Figure C2](#), [Figure C3](#), [Figure C4](#), and [Figure C5](#) in [Appendix C](#).

⁶⁷That is, they do not experience the same number of judicial mentions at the same times. Because of this, we group works into simplified treatment cohorts. See [Section 4.2](#) below.

learning, and LLM classification purposes. For each realist or originalist work, the JSTOR data includes matches to articles or books that cite that work. Each match contains information on the keywords used in the article that cite the realist or originalist work. These keywords summarize that citing article. Consider, for example, that one match citing Jerome Frank’s *Law and the Modern Mind* is [Rumble \(1965\)](#). The keywords for the Rumble article are:

legal realists; legal realism; llewellyn; karl llewellyn; jurisprudence; judicial; american legal; legal norms; ideals; judicial decisions

Each article that cites one of the realist or originalist works has its own keywords.⁶⁸

We use an LLM classifier to categorize citing academic articles into particular legal subjects. Using the JSTOR keywords, author name, and article title, we classify the articles which cite the set of realist and originalist works into legal subject categories. The instructions we gave to ChatGPT 4o-mini are as follows:

GPT prompt: “You are a legal scholar. Given the following article metadata, classify it into exactly one of the following categories: Property, Contracts, Torts, Constitutional Law, Administrative Law, Legislation, Corporate, Tax, Criminal Law, or Jurisprudence. Respond only with the category.”

Metadata provided:

- Title: [title of article]
- Author: [author name]
- Keywords: [keywords]

GPT output: Category: [category]

⁶⁸In [Figure A5](#), we show a word cloud of the 100 most common word stems among these keywords for both the realist and originalist corpora. The two word clouds contain 54 words in common. These are mostly words that relate to law and judging generally. The three most mentioned words in articles citing legal realist works that are also mentioned by originalist-citing articles are: law, legal, and judici. The three most mentioned word stems in articles citing originalist works that are also mentioned by realist-citing articles are: constitut, law, and judici. The three most mentioned words in articles citing legal realist works that are *not* mentioned by originalist-citing articles are: contract, realism, realist. The three most mentioned words in articles citing originalist works that are *not* mentioned by realist-citing articles are: origin, mean, and scalia. Unsurprisingly, the originalist-citing articles also include many keywords related to constitutional clauses that are not included in the top 100 realist keywords, such as: fourteenth, equal, protect, religion, and abort.

Sample output for two articles from each subject category are included below in [Table 2](#). The GPT-output categories seem reasonably accurate based on the titles and keywords. In [Table 3](#), we perform a heuristic accuracy assessment to assess the accuracy of the labels. Specifically, we look for whether the GPT-output category stem word appears in the article title or keywords. For example, using regular expressions (“regex”) we see whether the word “contract” appears in the article title or keywords columns for each JSTOR article labeled as a Contracts article. In the case of Contracts, 79.18% of articles labeled by GPT as a contracts article include the word “contract” in the title of the article or in the set of keywords.

Table 2. Sampled Articles by Legal Category

category	title	author
Administrative Law	Statutory Interpretation In The Administrative ...	Colin S. Diver
Administrative Law	Aliens: Deportation: Statutory Construction	William K. Jackson
Constitutional Law	Constructed Constraint And The Constitutional ...	Curtis A. Bradley; Neil S. Siegel
Constitutional Law	State Surveillance And The Right ...	Gautam Bhatia
Contracts	The Advent Of Recovery On ...	Kevin M. Teeven
Contracts	The Restatement Of The Law ...	Edwin W. Patterson
Corporate	Winner Of The SIs Annual ...	Dr Daniel Attenborough
Corporate	Behind The Cloak Of Corporate ...	Fenner L. Stewart
Criminal Law	Fear Of Forgiving: Rule And ...	Margaret Colgate Love
Criminal Law	Felony Murder And Mens Rea ...	Guyora Binder
Jurisprudence	Toward A Black Legal Scholarship: ...	Jerome Mccristal Culp, Jr.
Jurisprudence	The Problems Of A Functional ...	Felix S. Cohen
Legislation	The Interaction Of Legislation And ...	M. B. Hooker
Legislation	Statutory Language And The Purposive ...	Arthur S. Miller
Property	Whose Animals? A History Of ...	L. Naughton-Treves; Simon Batterbury; Tony Bebbington
Property	Effects Of Ownership Rights On ...	Hildy S. Ross
Tax	Recent Supreme Court Decisions In ...	Randolph E. Paul
Tax	The Role Of Prescription In ...	Anthony P. Polito
Torts	The Intellectual Origins Of Torts ...	G. Edward White
Torts	Autonomes Fahren - Bremsen Ethik ...	Johannes Schmidt; Matthias Klatt

Table 3. Articles with Category in Title or Keywords (Regex Check)

category	pct_matched
Administrative Law	49.23
Constitutional Law	66.14
Contracts	79.18
Corporate	63.55
Criminal Law	79.94
Jurisprudence	16.10
Legislation	58.09
Property	67.03
Tax	99.38
Torts	37.70

The classification results into legal subjects across the originalist and realist corpora are shown below in [Figure 4](#). The constitutional law and jurisprudence categories are most represented in each corpus. Beyond these top two categories for each corpus, the legal subject classifications of the citing articles aligns relatively with the foci of each of the legal movements. Because of its emphasis on the Supreme Court’s constitutional jurisprudence, it makes sense that articles citing originalist works deal more with administrative law and legislation than the traditional common law subjects. By contrast, articles which cite one of the realist articles are more focused on the first-year common law subjects: Contracts, Property, and Torts.

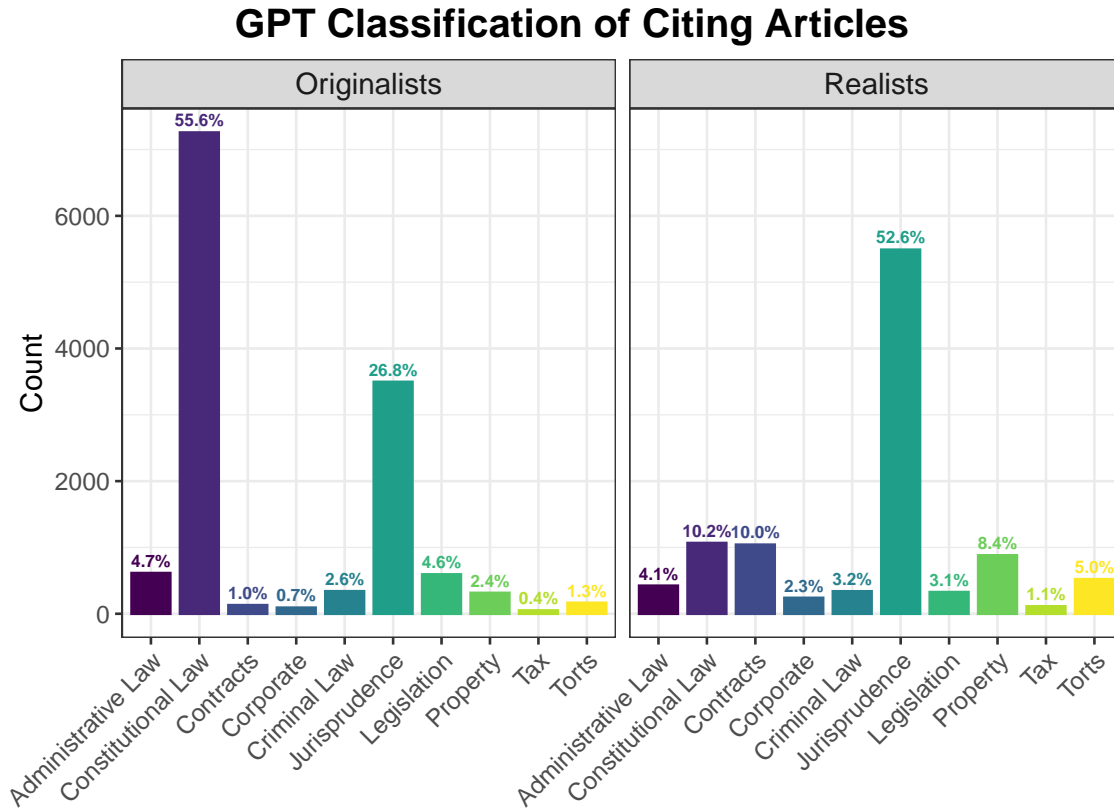


Figure 4. ChatGPT Legal Subject Classifications of Citing Articles

Note: The bars show the raw counts of citing articles under each legal subject. The percentage of articles in the corpus (realists or originalists) which fall into the legal subject are also listed at the top of each bar.

3.3.2 Citations & Mentions

The primary data of interest are the two time series for each work. Recall the two time series are: (1) the number of academic citations to an originalist or realist work (JSTOR Citations); and (2) the number of mentions for an originalist or realist work in published judicial opinions (Court Mentions).

In Figure 5, we show the basic linear fit for each author’s total JSTOR citations on the total court mentions. The totals are plotted on the log scale for visualization purposes. We label the names for authors with four or more works in each dataset. The linear fits of both the realists (yellow line) and originalists (purple line) have positive slopes. That is, the total number of academic JSTOR citations is increasing in the total number of court mentions.

Aggregated Jstor Citations & Court Mentions

Originalist and Realist Authors




Figure 5. Overview: Relationship between Court Mentions and Academic Citations

Note: Each dot represents a unique author grouping and has volume proportional to the number of total works. Authors with four or more total works are labeled.

Estimates calculated using regressions, an example of which is the line of best fit displayed above, do not necessarily imply a causal relationship. Since the linear fit is for the cumulative series, it does not respect the ordering of court mentions and academic citations, which is necessary for a causal implications.

The regression specification that produced the line of best fit above can be spelled out as follows:

$$\text{Total JSTOR Citations}_{ig} = \alpha_g + \beta_g \times \text{Total Court Mentions}_{ig} + \text{error}_{ig}$$

The subscripts g relate to the author group (realists or originalists). The independent variable $\text{Total Court Mentions}_{ig}$ is the total court mentions for author i in group g .⁶⁹ The goal is to obtain an estimate for β_g (the slope in the line of best fit). Such an estimate can be interpreted as the average increase in citations that result from one

⁶⁹Note also that the variables are on the log scale in the specifications.

more court mention for the author group.⁷⁰

Since this specification obscures important variation within publications in the author group and variation in the court mention path, a different regression that is at a more granular work and year level can be also specified. Now the i subscript refers to the published work, the subscript t to the year, and the subscript g to the group of authors (realists or originalists):

$$\text{Cumulative JSTOR Citations}_{igt} = \alpha_g + \beta_g \times \text{Cumulative Court Mentions}_{igt} + \text{error}_{igt}$$

All citations and mentions are calculated as cumulative series as in [Figure 1](#) and [Figure 3](#). Now the interpretation of β_g is the average increase in academic citations in a year for an additional court mention in that year (for group g).

Notice, however, that this specification assumes that an effect is instantaneous. Given the possible mechanisms by which judges learn of and mention legal scholarly works in their opinions, such an assumption is overly restrictive. As [Figure 2](#) illustrates, scholarly works are rarely mentioned immediately after publication. Therefore such an instantaneity assumption can induce a downward bias on the estimate β_g for each group of authors g .⁷¹

Regression evidence, moreover, may be particularly problematic for causal inference where, as here, there is sparsity in the treatments. That is, we only observe limited court mention paths out of a large space of conceivable such paths. Estimates such as β that attempt to encode this process then have limited information to “chew on.” To make matters worse, naive linear regression makes comparisons between units on different causal paths and does not necessarily have a causal interpretation.

The previous discussion implies that in this dynamic setting there may be several causal quantities of interest, some requiring stronger assumptions to recover than others. We now turn to detailing the challenges to causality in this setting and explaining the need for bounding.

⁷⁰While the linear fit seems appropriate to the realists, it cannot account for a cluster of originalist works that receive no court mentions but have academic citations that exceed author groups that received many court mentions.

⁷¹There are alternatives to account for the violation of instantaneity. But they require either culling the sample or making strong assumptions about the judicial process. For example, we may consider works only after their first judicial mention. Or we may assume that works take some number of years to be recognized by the judiciary after publication, and truncate both series to account for this lag. Both are imperfect solutions and likely to overestimate the effect of judicial acknowledgments because the information removed occurs where judicial mentions are small.

4 Methodology: Bounding Causal Effects

The canonical machinery of causal inference is derived under simple treatment paths. The familiar potential outcomes framework is introduced in a setting where each unit’s treatment path is only one time period long. There are then only two possible treatment paths: treated or untreated. It is assumed there are a sufficient number of treated and untreated units that can be used to impute the counterfactuals for each potential outcome. Under different assumptions, quantities like the average treatment effect, average treatment effect on the treated, or average treatment effect on the untreated can be estimated.

In dynamic treatment settings, treatments paths are many-periods long. The potential outcomes framework still applies but is more involved. More assumptions must be made, for example, about whether past treatments matter for later treatments, whether units dynamically select treatments (e.g. “dynamic selection on gains”), and whether units anticipate treatment. Yet, even in these dynamic settings, some treatment paths are simpler than others. If all treatment occurs at one period, the paths can essentially be reduced to the simpler setting by making simplifying assumptions.⁷²

Because the number of potential outcomes grows for each individual as treatment paths increase in complexity, we now turn to certain simplifying assumptions on how full treatment history affects potential outcomes, which reduces but does not eliminate the complexity. We then propose our Bounded Deviations strategy and show how these assumptions can point-identify or bound causal effects of interest in our setting.^{73 74}

This section first sets up the counterfactual framework for this setting and describe how it creates a dimensionality problem. Next, we briefly discuss how some simplifying assumptions can be difficult to motivate, hard to interpret, and result in biased causal estimates. Finally, we explain our counterfactual methodology, illustrating it with a simulated example and visualizations.

⁷²A well-studied dynamic setting is the staggered adoption setting where the treatment may occur at different times of the path, but any unit if treated at all is only treated once. In such a setting, treated units are usually assumed to retain memory of treatment forever. See, e.g., [Callaway and Sant’Anna \(2021\)](#). See also discussion in [Appendix B](#).

⁷³Our approach is a Manski-type partial identification assumption. See, e.g., [Manski and Pepper \(2018\)](#).

⁷⁴We compare our approach to the canonical difference-in-differences (DID) approach in [Appendix D](#).

4.1 Counterfactual Framework

We now introduce notation to describe the potential outcomes as a function of the treatment histories or treatment paths for each piece i in each time period t .⁷⁵ Associated with each i are treatment histories (i.e. judicial mention paths) $\vec{K}_i = (k_{i1}, \dots, k_{iT})$, where k_{it} is the observed treatment level for piece i at time period t .⁷⁶ So, for instance, the treatment history $(0, 1, 1)$ means a piece was mentioned zero times in the first period and once in the second and third periods. We imagine that any piece could have been subjected to any alternate judicial mention path and have shown a different trajectory of academic citations, which we refer to as a potential outcome.

We write the potential outcomes (i.e. potential academic citations) for each piece i in time period t under treatment path $\vec{k} = (k_1, \dots, k_T)$ as:

$$Y_{it}(\vec{k}) = Y_{it}(k_1, \dots, k_T)$$

This is the number of academic citations piece i would receive in time t if it experienced judicial mention history (k_1, k_2, \dots, k_T) .⁷⁷ Finally, we define the potential outcome matrix \mathbf{Y}_i for piece i . This matrix collects the many potential outcomes under counterfactual treatment histories.⁷⁸ For example, consider a two-time period model with a binary treatment possible in each period and an individual i that didn't receive treatment at time $t = 1$ but did receive treatment at time $t = 2$.⁷⁹ Then we write i 's potential outcome matrix \mathbf{Y}_i as follows:

$$\mathbf{Y}_i = \begin{bmatrix} Y_{i,t=1}(0,0) & Y_{i,t=2}(0,0) \\ Y_{i,t=1}(0,1) & Y_{i,t=2}(0,1) \\ Y_{i,t=1}(1,0) & Y_{i,t=2}(1,0) \\ Y_{i,t=1}(1,1) & Y_{i,t=2}(1,1) \end{bmatrix} = \begin{bmatrix} ? & ? \\ Y_{i,t=1}(0,1) & Y_{i,t=2}(0,1) \\ ? & ? \\ ? & ? \end{bmatrix}$$

The unobservable potential outcomes for i are denoted by the question marks. They define the **identification problem**: because unit i has does not receive treatment at

⁷⁵All time periods can be collected and ordered in a set $\mathcal{T} = \{1, \dots, T\}$

⁷⁶Note that, in principle, each k_{it} can be multidimensional if there are multiple treatments, but we focus on scalar k_{it} within a discrete set $\mathcal{K} = \{0, 1, \dots, K\}$.

⁷⁷Implicit in this notation is that there are no treatment “spillover” effects, i.e. the potential outcome for piece i at time t depends only on i 's treatment path.

⁷⁸For a fixed ordering of the potential treatment paths $\mathcal{K} = \{\vec{k}_1, \vec{k}_2, \dots\}$, the (i, j) -th entry of \mathbf{Y}_i is given by $Y_{i,t=j}(\vec{k}_i)$.

⁷⁹In other words, their observed treatment path is $\vec{K}_i = (k_{i1} = 0, k_{i2} = 1)$.

$t = 1$ and does receive treatment at $t = 2$,⁸⁰ the outcomes for i in period t at different treatment histories are unobservable.⁸¹

4.2 Curse of Dimensionality & Treatment Path Cohorts

The number of these unobservable potential outcomes grows as the number of time periods and number of treatments grow.⁸² This creates what is sometimes called “a curse of dimensionality” problem: allowing for i ’s potential outcome at time t to depend on the full treatment path increases the complexity of the identification problem.⁸³ In terms of our potential outcome matrix, the problem is that the number of question marks inside the matrix is growing too quickly.

To take control of the explosion of paths, we define four treatment cohorts, which simplifies the treatment path space. The first cohort is the never-treated cohort: these units never receive any judicial mentions. The remaining three cohorts are those that receive at least one judicial mention. They are split according to whether they receive a court mention Early, Mid, or Late in the movement. The years were determined to ensure a roughly equal split among the three treated cohorts. This resulted in the following cohort splits for each legal movement:

Realists		Originalists	
Cohort	Number Works	Cohort	Number Works
Early (1925–1945)	9	Early (1974-2009)	31
Mid (1946–1975)	7	Mid (2009-2019)	28
Late (1976–2025)	7	Late (2020-2025)	32
Never-Treated	8	Never-Treated	100

Table 4. Cohort Distribution for Realists and Originalists

In each table the years in parentheses indicate the times that the cohort first received a court mention. For example, if a realist work was first mentioned in a published court opinion between 1925–1945, then it is in the Early treated cohort.

⁸⁰That is, i has treatment history $\vec{K}_i = (0, 1)$.

⁸¹The unobservability of some counterfactual outcomes is often called the “fundamental problem of causal inference.” See [Holland \(1986\)](#).

⁸²As the cardinality of \mathcal{K} and \mathcal{T} increase, that is.

⁸³To address the dimensionality problem, researchers typically restrict the dependence of the potential outcomes $Y_{it}(\cdot)$ on the treatment history. We discuss further in [Appendix B](#) how different assumptions that are (implicitly or explicitly) used in practice can reduce the dimensionality of the identification problem.

The cohorts are roughly equal in terms of size, except that for the originalists there are many more Never-Treated units.

4.3 Bounded Deviations & Partial Identification

We now turn to explaining the framework we use to bound treatment effects in this setting. As in [Manski and Pepper \(2018\)](#), we use assumptions that bound causal effects for a pre-determined level of tolerance, technically known as a tuning parameter.⁸⁴ Our approach—which we call **Bounded Deviations**—uses information from units before their treatment paths diverge (i.e. *pre-divergence*) to impute counterfactual paths when the treatment paths no longer agree (i.e. *post-divergence*). Intuitively, the method imposes that the average treated group counterfactual must be “close” enough to the average control group observed outcome. There are many choices to encode this closeness. But we use the maximum observed pre-treatment deviation, which, to be more specific, is the largest difference in annual academic citation count between two series since they were both in publication.⁸⁵ Our method requires the treated group counterfactual to be within the maximum pre-divergence difference between observed treated and control average outcomes.⁸⁶ We use this method to estimate the average treatment effect on the treated (“ATT”), which is the average amount treatment changed outcomes for those actually treated.⁸⁷

For demonstration and visualization purposes, we first illustrate the Bounded Deviations approach with mock simulated data. In [Figure 6](#), we consider two average group outcomes for units on different treatment paths \vec{k}_1 and \vec{k}_2 . For simplicity, we call the \vec{k}_1 path “Treated” and the \vec{k}_2 path “Untreated.” These treatment paths are the same until 1950, after which they diverge. The observed average for Treated units is plotted as the blue line from 1900 to 2000.⁸⁸ We also plot the observed average for Untreated units in red only before 1950 for visual clarity.⁸⁹ The maximum pre-

⁸⁴Our approach is a variant of the “Bounded Variation” DID-style methodology in [Manski and Pepper \(2018\)](#), adapted to a setting with long treatment histories.

⁸⁵This choice was motivated by its ease of interpretation, but other choices include the average deviation before treatment divergence.

⁸⁶See [Appendix D](#) for full details.

⁸⁷For familiar readers, the Bounded Deviations method we use nests the widespread difference-in-differences method (DID) strategy. For a discussion of how the bounds from our method contain the usual DID estimates if parallel trends does hold, see [Appendix D](#). We also illustrate how causal conclusions are sensitive to the tuning parameter assumption in [Appendix E](#).

⁸⁸Formally this is the observed conditional expectation function $\mathbb{E}[Y_{it}(\vec{k}_1)|K_i = \vec{k}_1]$, which is plotted in blue.

⁸⁹Formally, this is the observed conditional expectation function $\mathbb{E}[Y_{it}(\vec{k}_2)|K_i = \vec{k}_2]$, which is shown

divergence deviation between the average for the two groups is 1.69 and occurs shortly after 1925.⁹⁰

The problem of causal inference here (i.e. the “identification problem”) is what counterfactual outcome would units on one path experience if they were instead on another treatment path? In the left panel of Figure 6, the blue-striped region is the counterfactual region for the blue Treated units if they had instead been Untreated.⁹¹ Recall our identifying assumption: if the two treatment paths had not diverged, average Treated outcomes would be “near” the observed average for Untreated units.⁹² How close would they be? They’d be within the maximum pre-divergence deviation.⁹³ In Figure 6, the maximum pre-divergence deviation is 1.69 and the blue-striped region is the counterfactual region under the assumption that blue units would’ve continued to be within 1.69 of the observed red unit outcomes.

in red.

⁹⁰Note that what is plotted in Figure 6 are the conditional expectation functions (CEFs) for units on a particular path. For example, the solid-red line is the average observed outcome for units on path k_2 in each time period.

⁹¹Recall that Treated units are those that actually experienced treatment path \vec{k}_1 . Untreated units are those that instead experienced treatment path \vec{k}_2 .

⁹²Mathematically, we are saying the CEF for units on treatment path \vec{k}_1 would be close to the CEF for units with path \vec{k}_2 .

⁹³When the multiple $C = 1$.

Bounded Deviation Identification (Fixed $C=1.0$)




Figure 6. Bounded Deviations Example

Note: In the left panel, the two lines are average outcomes for Treated (blue) and Untreated (red). The black-dotted vertical line is the path divergence time period. The blue-striped region is the counterfactual region for blue Treated units: this is where counterfactual outcomes for Treated units would be if they had been Untreated. The maximum pre-divergence deviation is 1.69 and is shown by the black vertical line around 1925. In the right panel, the y -axis is the average treatment effect on the treated, with a black-dotted horizontal line is plotted for zero effect. The upper and lower bounds on the ATT for each post-1950 time are plotted as dashed-yellow and solid-green lines, resp.

This counterfactual region implies bounds on treatment effects. These bounds are shown in the right panel of Figure 6. In the panel, the lower bound is plotted as a solid-green line and the upper bound is plotted as the dashed-yellow line. The dotted horizontal line at zero is also shown and corresponds to no effect (i.e. $ATT = 0$).

What do these bounds tell us? Before 1975, our bounds include the possibility of a zero effect—i.e. the dotted line at zero lies between the lower bound (solid-green line) and the upper bound (dashed-yellow line). Notice that this corresponds to the pre-1975 solid-blue line lying entirely within the blue-striped counterfactual region in the left panel. Moreover, we cannot say (identify) what the direction (sign) of the effect is. That is, we cannot know whether Treated units would have experienced

higher or lower outcomes relative to their counterfactual outcomes had they been Untreated.⁹⁴

After 1975, by contrast, the bounds do not contain the zero effect. That is, the dotted line lies below the lower bound on the *ATT* (solid-green line). Therefore, we can identify that the treatment effect of being Treated rather than Untreated increased the average outcomes for Treated units post-1975.⁹⁵ In other words, the treatment effect is sign-identified post-1975 because the upper and lower bounds are both positive, even though we cannot say precisely the magnitude of the effect.

The bounds we obtain are sensitive to a counterfactual tuning parameter C . The plots above assume $C = 1$, which intuitively means that the pre-divergence and post-divergence periods differ only in the judicial mentions but otherwise are “sufficiently alike.” How alike they are is governed by how large or small C is. Large values of C will increase the size of the counterfactual region and suggest that the past is a relatively weak guide to the future. By contrast, values of $C < 1$ imply that the counterfactual outcome is even closer to the observed outcome than it was pre-treatment divergence. We show how this affects the upper and lower bounds in [Appendix E](#).

5 Results: How Courts Impact Legal Movements

In this section, we use the methods developed above to estimate bounds on the causal effects of judges on legal realism and originalism. We quantify the impact of judicial mentions on the reception of academic legal ideas. Specifically, we bound the causal effects of judicial citation to certain legal realist or originalist works on the popularity of those works over time. We illustrate how to estimate upper and lower bounds for works if they had experienced alternative (i.e. counterfactual) treatment paths.

5.1 Individual Treatment Effects: An Illustration with Two Realist Works

First, we illustrate the methodology with a simple example of two contemporary realist works. Our causal strategy can be easily adapted to estimate individual (unit-

⁹⁴That is, we cannot identify the sign if those Treated units on path \vec{k}_1 had instead been on the “Untreated” path \vec{k}_2 for years prior to 1975.

⁹⁵Again, nothing here requires a binary treated-untreated setting: this methodology can be applied if some units are on treatment path \vec{k}_1 rather than another treatment path \vec{k}_2 .

level) treatment effects.^{96 97} We consider, for example, comparisons between two early realist works—from Holmes and Gray—in [Figure 7](#). In the left panel, the observed JSTOR citations for Holmes’ “The Path of the Law” and Gray’s *The Nature and Sources of the Law* are plotted in blue and red, respectively. For many years, neither work receives judicial mention according to the CourtListener data. After 1947, however, the treatment paths for the two works diverge, which is represented by the black dotted line. After that time, the two works receive a different number of court mentions. The levels of the court mentions are shown by the volume of the blue and red dots for each work. As shown in the left panel of [Figure 7](#), Holmes’ “The Path of the Law” receives many more judicial mentions (larger blue dots).

⁹⁶As in the synthetic control literature. See, e.g., [Abadie et al. \(2010\)](#) and [Abadie et al. \(2015\)](#).

⁹⁷In other words, we can relax the simplifying assumptions about the path space to bound unit-level causal effects. Recall as discussed above that few of the legal works we study experience the exact same treatment paths (judicial mentions) over the entire time. Instead, the different works are mentioned at different times and at different intensities (levels). For example, some legal works are mentioned early by many judges but then receive few (if any) judicial mentions in the late-20th or early-21st century. By contrast, some works only begin to receive a few judicial mentions many years after initial publication. And a select few works (e.g. *Corbin on Contracts*) have received high-levels of judicial mentions most years since publication. Our approach can allow for the potential popularity for each legal work to differ depending on when and how often judges mention the work in (published) court opinions. See further discussion in [Appendix B](#).

Unit-Level Bounded Deviations Identification




Figure 7. Individual (Unit-Level) Treatment Effects Illustration

Note: The left panel shows the observed academic citations for two works. The dots represent years in which the work received court mentions, with the volume of the dots corresponding to more mentions. The right panel shows the individual treatment effect upper and lower bounds for tuning parameter $C = 1$.

Using our Bounded Deviations approach for individual units, we can bound the causal effects from Holmes’ work experiencing a different treatment path than Gray’s work. The blue-striped region in the left panel is the counterfactual region, as discussed above: it is where we assume the citations for “The Path of the Law” would have been if it had stayed on the same judicial mention path as Gray’s work. In the right panel of [Figure 7](#), we use this counterfactual region to compute the upper and lower bound on the individual treatment effect. As you can see, the sign of the treatment effect in earlier years includes the zero effect possibility. However, in later years the treatment effect is positive. This is because the citations to Holmes’ work are much higher than the top part of the counterfactual region.⁹⁸

Although our method allows for bounding dynamic treatment effects for any pair, we recommend this only for those works that share a long enough pre-divergence period. Usually this means both pieces were in publication for several years before one was mentioned by the judiciary. Individual treatment bounds give us the most granular information, but the number of possible comparisons that could be made

⁹⁸See [Appendix B](#) for further discussion.

can be overwhelming. For example, even with 20 contemporary pieces there are 190 possible comparisons that we could make.⁹⁹

Since it would be difficult to summarize such a large number of comparisons, we coalesce our treatment paths to cohorts, and continue our analysis on these cohort-level series. Cohorts are defined based on when a legal piece received their first judicial mention.

5.2 Cohort Averages

Here, we turn to identifying effects for entire cohorts of realist and originalist works. We first display the cohort average JSTOR citations for the realists and originalists. Recall that our cohorts are defined as follows: the Never-Treated cohort never receive any judicial mentions, the Early treated cohort first receives at least one judicial mention 1925–1945, the Mid treated cohort first receives at least one judicial mention 1946–1975, and the Late treated cohort first receives at least one judicial mention 1976–2025. The average number of yearly JSTOR citations for each cohort and corpus (realists and originalists) are shown below in [Figure 8](#).

⁹⁹In between calculating the individual treatment effect for any pair and a cohort averages is the possibility of calculating an individual treatment effect of a select piece against the average of all other never mentioned units, such as in a synthetic control, set-up. But even with this reduced number of comparisons we are left with too many comparisons to make any informative statements about the impact of judicial mentions.

Cohort Average Citations




Figure 8

Note: The divergence periods that demarcate the treatment cohorts are plotted as vertical black dotted lines.

For both originalists and realists, the never-treated averages are, in most time periods, lower than the three treatment cohort averages. In general, the Early treated average (in purple) is higher than the yearly citations received for other cohorts after Early Divergence. Before Early Divergence, however, observe that the Early, Mid, and Late cohorts receive a similar number of yearly citations on average.

These cohort average citation series are the basic “building blocks” that we use to compute causal effects. The results that follow use pre-divergence deviations between the series and our counterfactual tuning parameter C to compute upper and lower bounds on the treatment effect of a judicial mention.¹⁰⁰

5.3 The Effects of Court Mentions on Legal Realism

In our first application, we estimate the effect of judicial mentions on the reception for the realist thinkers mentioned above. For ease-of-reference, the cohorts for comparison purposes are as follows: the Never-Treated cohort, which received zero court mentions;

¹⁰⁰The Early treated cohort has a very short (case of legal realism) or no (case of originalism) pre-divergence time overlap with other cohorts. This makes it difficult or impossible, resp., to make appropriate comparisons between the early-mentioned and never-mentioned cohorts, so we report only the comparisons between Mid vs Never-Treated and Late vs Never-Treated within the text. See [Appendix F](#) for further results.

works in the Early treated cohort, which first received court mention in 1925–1945; works in the Mid treated cohort, which first received court mention in 1946–1975; and works in the Late treated cohort, which first received court mention in 1976–2025.

Our first set of results estimates the effect from being in the Mid treated cohort versus having never been treated. In [Figure 9](#), we plot the upper bound (dashed-yellow) and lower bound (solid-green) on the ATT for being in the Mid treated cohort versus the Never-Treated cohort for units actually in the Mid treated cohort.¹⁰¹ The different panels are for different values of the identification tuning parameter C .¹⁰²

Realists: Mid vs Never-Treated Cohorts




Figure 9. Realists: Mid vs Never-Treated Cohorts

Note: In each panel, the black dotted horizontal line is plotted for $ATT = 0$, the ATT upper bound is the dashed-yellow line, and the ATT lower bound is the solid-green line. The years for which the upper and lower bounds are both positive are shaded green. The years for which the upper and lower bounds are both negative are shaded red.

The light green-shaded region in each panel are those years for which we identify the sign of the ATT to be positive. As shown in the left-most panel of [Figure 9](#), many of the 1980–2000 years sign-identify the ATT when $C = 1$. This means that the Mid units would have been much “closer” to the Never-Treated units if they had never received judicial mentions. Instead, the judicial mentions increased the popularity

¹⁰¹In the appendices, we show results for comparisons involving the Mid treated cohort when Arthur Corbin is omitted. Although Corbin is often associated with realism, *Corbin on Contracts* is, in certain respects, distinct from the other academic realist texts insofar as its focus is largely doctrinal. In the appendix, we show that the results are largely unchanged if Corbin’s work is omitted from the Mid treated cohort comparisons. See [Appendix F](#).

¹⁰²See [Appendix E](#) for further details.

(academic JSTOR citations) for realist works that were first treated in the 1946–1975 years (Mid cohort).

The effect is even more pronounced for realist works in the Late treated cohort. Relative to Never-Treated units, realist works that were first mentioned in judicial opinions from 1976–2025 received a positive boost in academic citations. This effect holds across values of the identification tuning parameter C , as shown below in Figure 10. In each panel, there are many years for which we sign-identify a positive ATT for Late treated units relative to the Never-Treated cohort. This is indicated by the large number of years which are shaded light green in each panel of Figure 10.

Realists: Late vs Never-Treated Cohorts

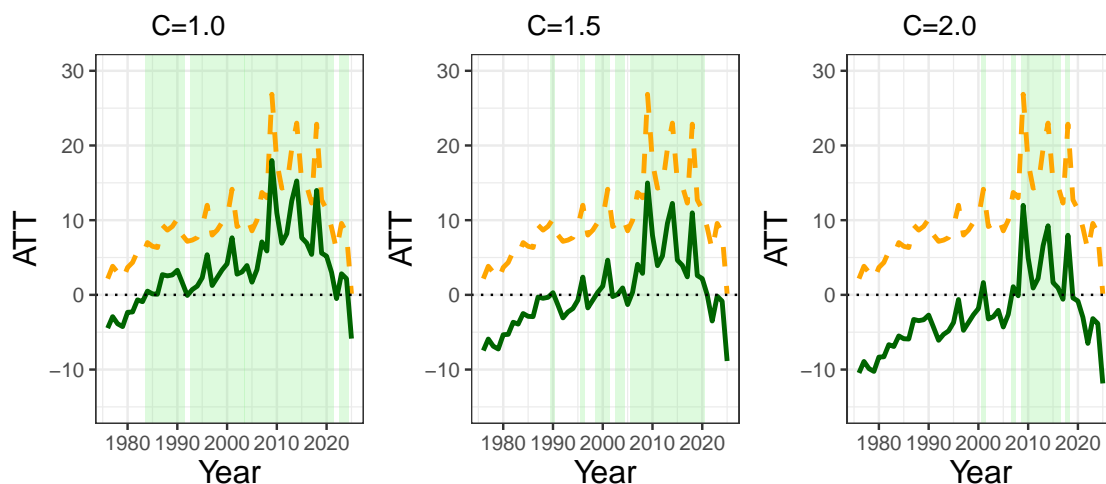


Figure 10. Realists: Late vs Never-Treated Cohorts

Note: In each panel, the black dotted horizontal line is plotted for $ATT = 0$, the ATT upper bound is the dashed-yellow line, and the ATT lower bound is the solid-green line. The years for which the upper and lower bounds are both positive are shaded green. The years for which the upper and lower bounds are both negative are shaded red.

These findings highlight an important difference between legal scholarship and judicial opinions. Existing empirical studies show that the value of court opinions depreciate rapidly. For example, Black and Spriggs II (2013) find that case precedent is less likely to be cited over time: new cases are about 30% likely to be cited, whereas 20 year old precedent is less than 10% likely to be cited.¹⁰³ By contrast, our results suggest some legal scholarship receives more citations later in life.¹⁰⁴ In other words,

¹⁰³See Figure 1 in Black and Spriggs II (2013), which shows that the predicted probability of citation as a function of the age of precedent decreases rapidly.

¹⁰⁴See also Figure 8.

court opinions age like milk but legal scholarship ages like wine.^{105 106}

Further comparisons between treatment cohorts are explored in [Appendix F](#). To summarize the results there, Early treated units received a positive boost from court mentions relative to Mid treated units, although this effect only appears in the 2000–2025 years. It also appears that being treated in the Mid cohort decreased citations relative to being treated Late. And, finally, being treated Early versus Late resulted in a positive effect for a few years in the 21st century, although most years the sign of the effect is not identified.

5.4 The Effects of Court Mentions on Originalism

In our second application, we estimate the effect of judicial mentions on the reception for the originalist thinkers mentioned above. For ease-of-reference, the originalist cohorts are: the Never-Treated cohort, which received zero court mentions; works in the Early treated cohort, which first received court mention in 1974–2009; works in the Mid treated cohort, which first received court mention in 2009–2019; and works in the Late treated cohort, which first received court mention in 2020–2025.

Our first originalist comparison is the Mid treated cohort versus the Never-Treated cohort. As shown in [Figure 11](#), the upper and lower bounds on the *ATT* for Mid treated units are shown in dashed-yellow and solid-green, respectively. For originalism, the sign of the *ATT* is identified for only one year and for only small values of the tuning parameter C . Moreover, the magnitude of possible *ATT* values is much smaller: the y -axis range is much smaller than in the realist results above. It is difficult to compare these results to the realist case because we have only observed the originalist movement for a much shorter span of time than the realism. Therefore, we caution against drawing a conclusion such as that judicial attention did not play a role in originalism since the effect may yet to be observed in data.

¹⁰⁵[Merritt and Putnam \(1995\)](#) uses Ronald Coase’s work on *Social Costs* as a specific example of scholarly writing that came to influence judicial decisions belatedly.

¹⁰⁶Cf. [Traynor \(1962\)](#) at 10. (“Time is with the law reviews. An age that churns up problems more rapidly than we can solve them needs such fiercely independent problem-solvers preoccupied with long-range solutions.”)

Originalists: Mid vs Never-Treated Cohorts




Figure 11. Realists: Mid vs Never-Treated Cohorts

Note: In each panel, the black dotted horizontal line is plotted for $ATT = 0$, the ATT upper bound is the dashed-yellow line, and the ATT lower bound is the solid-green line. The years for which the upper and lower bounds are both positive are shaded green. The years for which the upper and lower bounds are both negative are shaded red.

The comparison between Late treated originalist works and Never-Treated originalist works is similar. In [Figure 12](#), we do not identify the sign of the effect from being in the Late treated cohort rather than the Never-Treated cohort. This means we cannot say whether judicial mention increased or decreased the academic citations for Late treated originalist works relative to never being mentioned in court opinions. We do show, however, that the magnitude of the effect is small either way: for $C = 1.0$ in the left-most panel, for example, the average effect of a Late judicial mention is less than 5 in absolute terms. This means, for example, that the data is consistent with judicial mention increasing (or decreasing) the academic citations to Late treated originalist works by some amount less than 5.

Originalists: Late vs Never-Treated Cohorts




Figure 12. Originalists: Late vs Never-Treated Cohorts

Note: In each panel, the black dotted horizontal line is plotted for $ATT = 0$, the ATT upper bound is the dashed-yellow line, and the ATT lower bound is the solid-green line. The years for which the upper and lower bounds are both positive are shaded green. The years for which the upper and lower bounds are both negative are shaded red.

Again, we emphasize that we are working with data for only a small number of years.¹⁰⁷ Given the realist results, it appears that the impact of judicial mentions on academic citations takes many years (sometimes many decades) to noticeably affect academic citations. This makes sense given the time it takes for academics to appreciate particular works, process judicial changes to the law, and determine the canonical or important works in a particular legal tradition. In this case especially, originalist changes to constitutional law are active and ongoing. The effects of judicial originalism on legal academics may therefore take more time to appear in the data.¹⁰⁸

6 Conclusion

Existing empirical scholarship has primarily focused on how judges *use* academic legal works but has not estimated the *effect* of judges on legal academia. The primary intent of our work is to bring tools in the causal inference revolution to legal questions. Specifically, we show how a bounding approach can estimate the impact of

¹⁰⁷Particularly for the Late vs Never-Treated cohorts.

¹⁰⁸As in the realist case, see [Appendix F](#) for an additional result plot. There we compare Mid and Late treated originalist works and similarly find no sign-identification for C -values of 1.0, 1.5, 2.0.

judicial mention on citations to realist and originalist works. Estimating counterfactuals in settings where potential outcomes may depend on the entire history or path of treatment requires several simplifying assumptions. Bounding approaches can accommodate these simplifications by being more cautious about the causal range. As discussed above, the general strategy is to assume how unobserved counterfactual quantities evolve to determine the causal effects of treatment at different times (i.e. on cohorts). This approach can be applied to further legal settings with complicated dynamics.¹⁰⁹

This paper is relevant to the alleged divergence between the legal academy and the judiciary. In contrast to empirical analyses that test for a divergence by examining the entire universe of legal scholarship, we narrow our focus to two particular jurisprudential movements. Despite this restriction, our data shows large variation in when (and to what degree) judges mention scholarly work. Some works are mentioned early and often, whereas others are mentioned (if at all) long after publication. It is this variation that we exploit for causal identification to test whether judicial mentions matter for academics scholarship or not.

In our application, we find that courts impact the popularity of legal theories. Although it is too soon to say with certainty how judges affect originalist citations, we find positive effects of judicial mentions on citations to legal realist works relative to works that are never mentioned by judges. In the case of legal realists, we also find support for the conjecture that “time is on the side of law reviews,”¹¹⁰ in that delayed recognition by the judiciary does not imply that legal scholarship and judging are mutually irrelevant.

Because of our focus on the two movements, we cannot extrapolate our findings to all areas of the law. A challenge with doing the same analysis at the level of the entire scholarly corpus is to keep our comparison sets similar in their content. For we cannot compare scholarship about matters of constitutional interpretation to criminal procedure. Among other considerations, their importance to judging differs in the volume of such cases heard and the marginal impact of work across fields is non-uniform. Future empirical work must attempt to overcome these challenges to better understand the relationship between the courts and the academy.

¹⁰⁹As shown above, boundedness assumptions nest the standard parallel trends assumption and are robust when counterfactual trends are not parallel. This is important because much empirical legal work relies on the difference-in-difference strategy to infer causal effects. Violations of the assumptions underpinning that strategy can result in “heroic” estimates.

¹¹⁰Traynor (1962).

References

- Abadie, A., A. Diamond, and J. Hainmueller (2010). Synthetic control methods for comparative case studies: Estimating the effect of california’s tobacco control program. *Journal of the American Statistical Association* 105(490), 493–505.
- Abadie, A., A. Diamond, and J. Hainmueller (2015). Comparative politics and the synthetic control method. *American Journal of Political Science* 59(2), 495–510.
- Baude, W. (2015). Is originalism our law? *Columbia Law Review* 115, 2349–2408.
- Black, R. C. and J. F. Spriggs II (2013). The citation and depreciation of u.s. supreme court precedent. *Journal of Empirical Legal Studies*.
- Callaway, B. and P. H. Sant’Anna (2021). Difference-in-differences with multiple time periods. *Journal of Econometrics* 225(2), 200–230.
- Chilton, A. S. and E. A. Posner (2015). An empirical study of political bias in legal scholarship. *Journal of Legal Studies* 44.
- D’Amato, A. (1989). Can any legal theory constrain any judicial decision? *University of Miami Law Review* 43, 513.
- D’Amato, A. (1999). The effect of legal theories on judicial decisions. *Chicago Kent Law Review* 74, 517–527.
- Detweiler, B. T. (2020). May it please the court: A longitudinal study of judicial citation to academic legal periodicals. *Legal Ref. Servs. Q.* 39.
- Diamond, S. S. and P. Mueller (2010). Empirical legal scholarship in law reviews. *Annual Review of Law and Social Science* 6, 581–599.
- Edwards, H. T. (1992). The growing disjunction between legal education and the legal profession. *Michigan Law Review* 91(1), 34–78.
- Feldman, A. (2018). Empirical scotus: With a little help from academic scholarship. *SCOTUS Blog*.
- Fitzpatrick, B. T. and P. K. Varghese (2017). Scalia in the casebooks. *The University of Chicago Law Review* 84, 2231–42.
- Fronk, C. R. (2010). The cost of judicial citation: An empirical investigation of citation practices in the federal appellate courts. *Journal of Law, Technology, and Policy*.
- Goldsmith-Pinkham, P. (2024). Tracking the credibility revolution across fields. *Arxiv*.

- Goodman-Bacon, A. (2021). Difference-in-differences with variation in treatment timing. *Journal of Econometrics* 225(2), 254–277.
- Greene, J. (2009a). On the origins of originalism. *Texas Law Review* 88(1).
- Greene, J. (2009b). Selling originalism. *Georgetown Law Journal* 97.
- Gulati, M. and V. Sanchez (2002). Giants in a world of pygmies? testing the superstar hypothesis with judicial opinions in casebooks. *Iowa Law Review* 87, 1141.
- Ho, D. E. (2005). Why affirmative action does not cause black students to fail the bar. *The Yale Law Journal* 114, 1997.
- Holland, P. W. (1986). Statistics and causal inference. *Journal of the American Statistical Association* 81(396), 945–960.
- Kennedy, D. and W. W. Fisher (2006). *The Canon of American Legal Thought*. Princeton University Press.
- Leiter, B. (2005). *The Blackwell Guide to the Philosophy of Law and Legal Theory*. Wiley.
- Leiter, B. (2010). *American Legal Realism*. Wiley.
- Manski, C. F. (1993). Identification problems in the social sciences. *Sociological methodology*, 1–56.
- Manski, C. F. (1997). Monotone treatment response. *Econometrica* 65(6), 1311–1334.
- Manski, C. F. (2007). *Identification for Prediction and Decision*, Chapter 9. Harvard University Press.
- Manski, C. F. and J. V. Pepper (2018). How do right-to-carry laws affect crime rates? coping with ambiguity using bounded-variation assumptions. *The Review of Economics and Statistics* 100(2), 232–244.
- Merritt, D. J. and M. Putnam (1995). Judges and scholars: do courts and scholarly journals cite the same law review articles. *Chi.-Kent L. Rev.* 71, 871.
- Newton, B. (2012). Law review scholarship in the eyes of the twenty-first century supreme court justices: An empirical analysis. *Drexel Law Review* 4, 399–416.
- Posner, R. A. (1992). The deprofessionalization of legal teaching and scholarship. *Michigan Law Review* 91, 1921.
- Posner, R. A. (2016). *Divergent Paths: The Academy and the Judiciary*. Harvard University Press.
- Pound, R. (1908). Mechanical jurisprudence. *Columbia Law Review* 8(8), 605–623.

- Rambachan, A. and J. Roth (2023). A more credible approach to parallel trends. *Review of Economic Studies*, 2555–2591.
- Rumble, W. E. (1965). The paradox of american legal realism. *Ethics* 75(3), 166–178.
- Sawyer III, L. E. (2018). Principle and politics in the new history of originalism. *American Journal of Legal History* 57(2), 198–222.
- Sirico Jr., L. J. and B. A. Drew (1991). The citing of law reviews by the united states courts of appeals: An empirical analysis. *University of Miami Law Review* 45, 1051–1060.
- Tamanaha, B. Z. (2008). Understanding legal realism. *Texas Law Review* 87, 731.
- Thalke, R., E. H. Stiglitz, D. Mimno, and M. Wilkens (2023). Modeling legal reasoning: Lm annotation at the edge of human agreement. *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 9252–9265.
- Traynor, R. J. (1962). To the right honorable law reviews. *UCLA L. Rev.* 10, 3.
- Whittington, K. (2004). The new originalism. *Georgetown Journal of Law and Public Policy* 2(2), 599–613.


APPENDICES

A Appendix A: Additional Descriptive Plots

Appendix A includes additional descriptive plots. We include the following plots: two additional academic citation series for the legal realist works (cumulative vs yearly citations); court mentions for realist works yearly; and court mentions for each author aggregated across works and for each decade.

Cumulative Citations

By Legal Realists

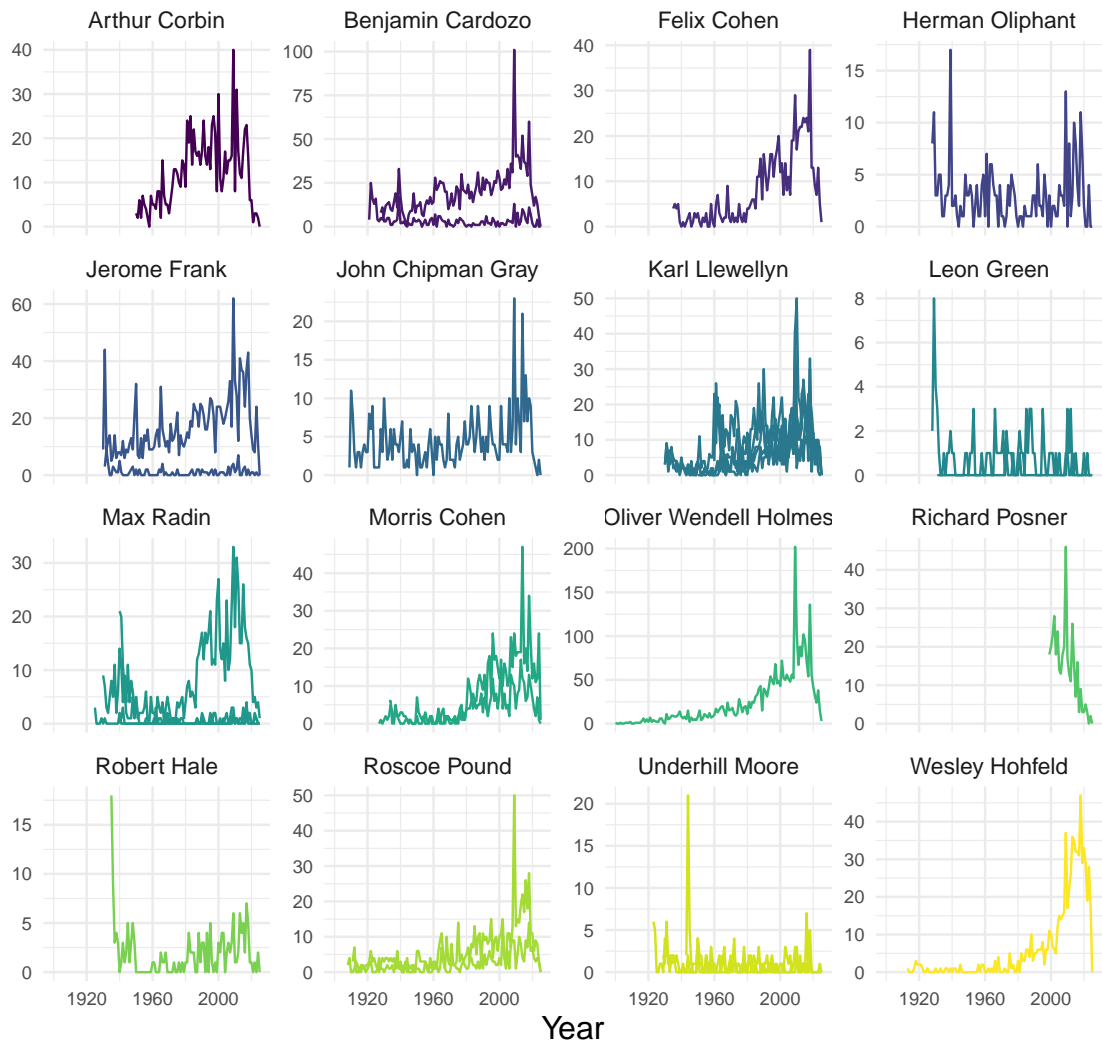


Appendix Figure A1. Cumulative Citations

Note: The y -axis differs for each legal realist. The different lines for each legal realist are for their different texts.

Yearly Citations

By Legal Realists



Appendix Figure A2. Yearly Citations

Yearly CourtListener Court Mentions By Legal Realists

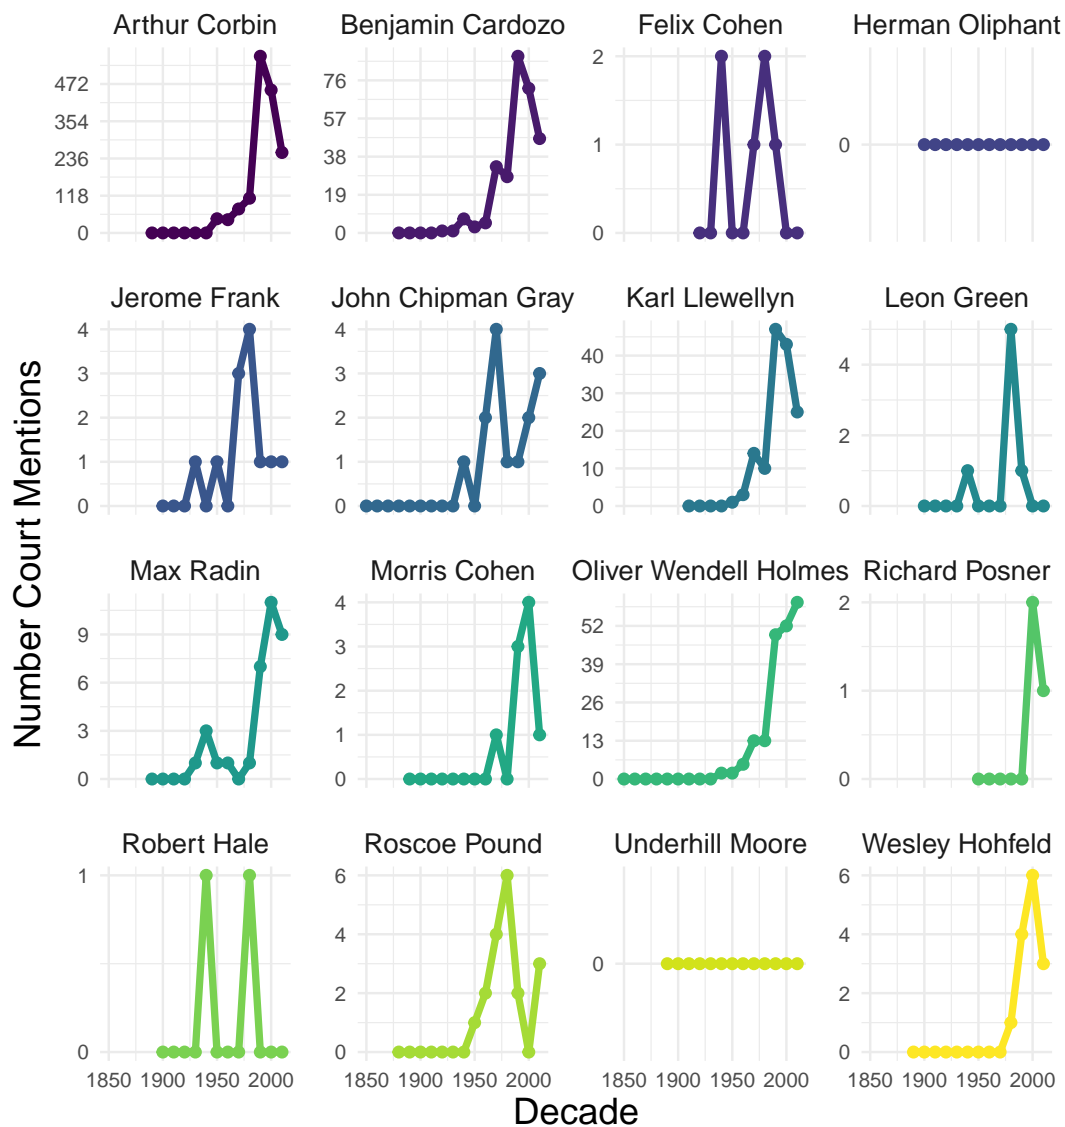


Appendix Figure A3. Yearly CourtListener Court Mentions

Note: The *y*-axis differs for each legal realist.

Court Mentions of Legal Realist Works

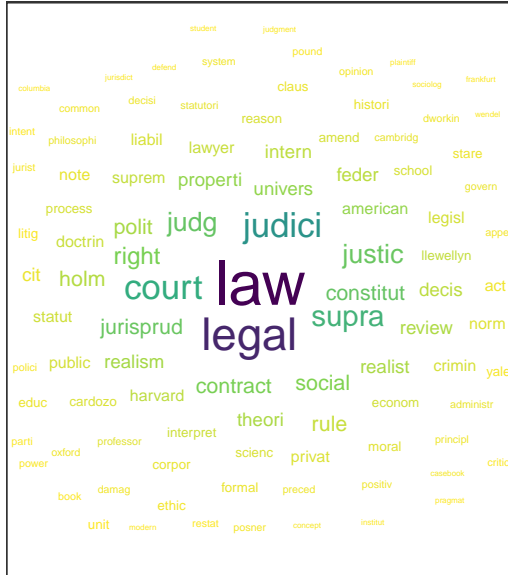
Total Mentions Across Works each Decade



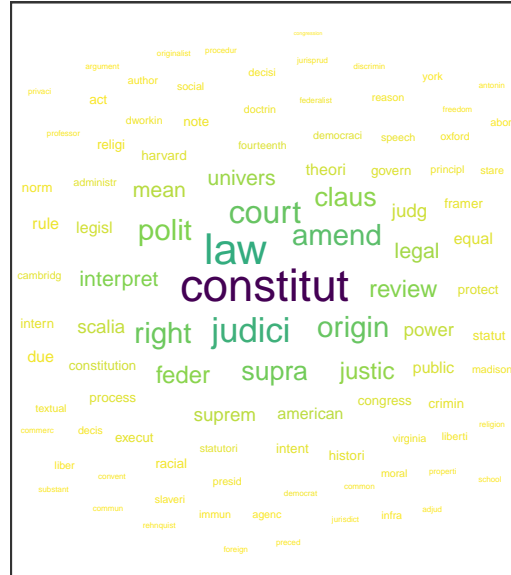
Appendix Figure A4. Legal Realist Court Mentions

Word Clouds by Legal Movement

Legal Realists: Top 100 Terms



Originalists: Top 100 Terms



Appendix Figure A5. Word Clouds by Legal Movement

B Appendix B: Counterfactual Complexity & The Curse of Dimensionality

Appendix B includes discussion and a plot of how potential outcomes grow as a function of the complexity of treatment paths. Researchers typically rely on simplifying assumptions rather than allowing for full path heterogeneity. Here, we discuss two types of assumptions that are (implicitly or explicitly) used in practice: memory and anticipation assumptions.

The first set of assumptions we term memory assumptions, which restrict how the potential outcome in period t can depend on treatments received prior to time t . Formally, we define *memory of degree m* to be the assumption that, for all i and in each time t , we have:

$$Y_{it}(\vec{k}) = Y_{it}(k_{t-m}, \dots, k_t, \dots, k_T)$$


This restricts the potential outcomes in period t to depend only on the treatment path up to m periods prior to t . For example, memory of degree zero (*no memory*) means that potential outcomes in time t depend only on treatments received during or after time t .

The second set of assumptions are anticipation assumptions. Commonly used the DID literature, this assumption restricts how potential outcomes in time t can depend on treatments received after time t . Formally, we define *anticipation of degree a* to be the assumption that, for all i and in each time t , we have:

$$Y_{it}(\vec{k}) = Y_{it}(k_1, \dots, k_t, \dots, k_{t+a})$$

This restricts the potential outcomes in period t to depend only on the treatment path up to a periods after t . For example, anticipation of degree zero (*no anticipation*) means that potential outcomes in time t depend only on treatments received during or before time t .

Figure B1 illustrates the dimensionality problem. We plot the number of potential outcomes under different memory assumptions and no anticipation. We consider settings with ten time periods and between 1-10 treatment dosages other than the baseline “no treatment” dosage. As seen in Figure B1, stronger memory assumptions greatly reduce the dimensionality of the problem. However, as the cardinality of the treatment set \mathcal{K} grows, the total number of potential outcomes remains large: under no memory and no anticipation, there are 100 total potential outcomes for each unit.



Appendix Figure B1. Potential Outcomes: Combinatorial Explosion

Note: This plot shows the total number of potential outcomes as a function of time periods (x -axis) and the number of treatments (line color). The number of treatments is the number of treatments in addition to a baseline “no treatment” dosage. The yellow line (1 treatment) is, e.g., a binary treatment in each period.

Even with these restrictions on the nature of the treatment path, counterfactual imputation is not as straightforward as the simpler setting. For one, there are now more counterfactual paths. Researchers might ask what the unit’s outcome would be if it were not treated at all, if it were treated earlier, or if it were treated later. The standard approach—comparing treated units to not-yet-treated units to obtain treatment effects—mimics the simpler settings described above.

These complications in dynamic settings are exacerbated as the complexity of the treatment path increases. Imputing counterfactuals with complex treatment paths is challenging for two reasons. The first is sparsity in the path space itself. Indeed, as shown in Figure 2, there is limited overlap in the treatment histories for units that are mentioned at least once in judicial opinions. The second is that even if many paths are observed, there will be too few observed units on each path to feasibly point identify a causal estimate. Both these make the usual machinery employed in causal inference difficult to work in these settings. Under such complex paths, there

is no clear demarcation between treatment and control units, particularly if all units get some treatment at some point.

Because it works to impute unit-level counterfactuals, the Bounded Deviations method we use is one possible solution. Our identification strategy partially identifies counterfactual paths for units receiving different treatment paths. If the length of a treatment path grows, the number of such counterfactual paths can be unmanageable quickly. To make this manageable, we could only limit to paths that we observe, or make parametric assumptions. For example, we could define a distance metric on treatment paths and assume that counterfactual outcomes are a function of the distance between paths. Counterfactual paths that are closer to the observed treatment path, e.g., may only have small, predictable deviations from the observed outcomes.

B.1 Possible Problems with Dimension-Reduction Assumptions

Memory-anticipation assumptions clearly reduce the dimensionality of the identification problem. Still, computational and interpretational problems may remain, especially in settings with numerous treatment levels (i.e. doses). But the main problem is that such assumptions can introduce bias and are difficult to motivate or justify in settings without well-founded theory. That is, if there are meaningful differences in potential outcomes along different treatment paths, comparing units along different paths can bias causal estimates.

Consider, for example, the simple two-period model with binary treatments available both time periods. If there is genuine anticipation, then the counterfactual averages along paths $(0, 0)$ and $(0, 1)$, say, may differ. That is, units that are untreated at time $t = 1$ but treated in time $t = 2$ may anticipate receiving treatment, act on it, and thereby experience different outcomes than units that are never treated (i.e. whose treatment paths are $(0, 0)$). If we then assume there is no anticipation, i.e. that $Y(0, 0) = Y(0, 1) \equiv Y(0)$, we improperly “mix” these units together. This can result in estimates of $\mathbb{E}[Y(0)]$ that differ substantially from the true population counterfactual averages $\mathbb{E}[Y(0, 0)]$ and $\mathbb{E}[Y(0, 1)]$. The direction and magnitude of the bias from assuming no anticipation can be difficult to predict or correct *ex post*.¹¹¹

In our setting, modeling anticipation may be especially difficult. Although researchers studying, e.g., the effects of particular pre-announced, well-discussed congressional or executive legal actions may be able to model anticipation, it is not clear

¹¹¹This problem is analogous to the problem of “forbidden comparisons” in staggered adoption settings. See, e.g., [Goodman-Bacon \(2021\)](#); see also [Callaway and Sant’Anna \(2021\)](#).

what legal academics anticipate in terms of judicial citations. It may be very difficult to anticipate which academic articles judges will cite, yet some law review pieces are prepared with particular court cases (esp. Supreme Court cases) in mind. Indeed, some legal academics may expect their writings to be successful because of pending judicial decisions and even file *amicus* briefs encouraging judges to engage with their scholarship. Particularly for a legal movement like the originalism movement in constitutional law, judges seem to be openly inviting academic discussion of particular legal issues.¹¹²

Simplifying the treatment path space is, therefore, not without possible downsides. However, in spaces with high-dimensional (time-varying) treatments, simplifying from full-path treatment heterogeneity to a few treatment cohorts is necessary to obtain overlap. As stated above, the flexible partial identification assumptions we propose above can work to impute unit-level counterfactuals (as in, e.g., synthetic control). This is because the unobserved potential outcomes depend on the observable potential outcomes but differences are “bounded.”

B.2 Individual Treatment Effects & Identifying Assumptions

Finally, we briefly sketch how one could consider individual treatment effect estimands. As in the synthetic control literature,¹¹³ we consider individual treatment effect estimands of the following form:

$$\Delta_i(t, \vec{k}, \vec{k}') = Y_{it}(\vec{k}) - Y_{it}(\vec{k}')$$

In what follows, we can obtain similar bounds on these individual treatment effects by the modified assumption below:

Assumption 1 (Individual Bounded Deviation). *Let $T_{div} = \operatorname{argmin}_t (k_t \neq k'_t)$ for two units i and j on distinct treatment paths \vec{k} and \vec{k}' , respectively. The following holds for each $t \geq T_{div}$ and some $C \in \mathbb{R}_+$:*

$$\left| Y_{it}(\vec{k}') - Y_{jt}(\vec{k}') \right| \leq C \cdot \max_{t < T_{div}} |Y_{it} - Y_{jt}|$$


¹¹²See, e.g., *Franklin v. New York*, 604 U.S. (2025) (Alito, J., concurring) (discussing the possibility of overturning Confrontation Clause precedent based on “[h]istorical research”).

¹¹³See, e.g., Abadie et al. (2010) and Abadie et al. (2015).

C Appendix C: Full Treatment Histories


Appendix C contains the full treatment histories for each of the realist and originalist works studied here. For each work, we plot the number of judicial mentions the work receives each year following publication. The levels are colored according to the legend at the top of each figure.

Court Mentions for Legal Realist Works




Appendix Figure C1. Legal Realists: Court Mentions

Court Mentions for Originalist Works (Part 1)




Appendix Figure C2. Originalists: Court Mentions (Part 1)

Court Mentions for Originalist Works (Part 2)




Appendix Figure C3. Originalists: Court Mentions (Part 2)

Court Mentions for Originalist Works (Part 3)



Appendix Figure C4. Originalists: Court Mentions (Part 3)

Court Mentions for Originalist Works (Part 4)



Appendix Figure C5. Originalists: Court Mentions (Part 4)

D Appendix D

Appendix D further describes our counterfactual framework and identification strategy.

D.1 Counterfactual Framework

We first introduce notation to describe the potential outcomes as a function of the treatment histories or treatment paths for each individual i in each time period t . Consider a setting with time periods $t \in \mathcal{T} = \{1, \dots, T\}$ and treatment histories $\vec{K}_i = (k_{i1}, \dots, k_{iT})$, where k_{it} is the observed treatment level for individual i at time period t . Note that, in principle, each k_{it} can be multidimensional if there are multiple treatments, but we focus on scalar k_{it} within a discrete set $\mathcal{K} = \{0, 1, \dots, K\}$.

We write the potential outcomes for each individual i in time period t under treatment path $\vec{k} = (k_1, \dots, k_T)$ as

$$Y_{it}(\vec{k}) = Y_{it}(k_1, \dots, k_T)$$

Implicit in this notation is that there are no treatment “spillover” effects, i.e. the potential outcome for unit i at time t depends only on i ’s treatment path. Finally, we define the potential outcome matrix \mathbf{Y}_i for individual i . For a fixed ordering of the potential treatment paths $\mathcal{K} = \{\vec{k}_1, \vec{k}_2, \dots\}$, the (i, j) -th entry of \mathbf{Y}_i is given by $Y_{i,t=j}(\vec{k}_i)$. For example, consider a two-time period model with a binary treatment possible in each period and an individual i with observed treatment path $\vec{K}_i = (k_{i1} = 0, k_{i2} = 1)$. Then we write i ’s potential outcome matrix \mathbf{Y}_i as follows:

$$\mathbf{Y}_i = \begin{bmatrix} Y_{i,t=1}(0,0) & Y_{i,t=2}(0,0) \\ Y_{i,t=1}(0,1) & Y_{i,t=2}(0,1) \\ Y_{i,t=1}(1,0) & Y_{i,t=2}(1,0) \\ Y_{i,t=1}(1,1) & Y_{i,t=2}(1,1) \end{bmatrix} = \begin{bmatrix} ? & ? \\ Y_{i,t=1}(0,1) & Y_{i,t=2}(0,1) \\ ? & ? \\ ? & ? \end{bmatrix}$$

The unobservable potential outcomes for i are denoted by the question marks. They define the **identification problem**: because unit i has treatment history $\vec{K}_i = (0, 1)$, the outcomes for i in period t at different treatment histories are unobservable.

D.2 Bounded Deviations: Identification

To begin, we specify the treatment effect estimand of interest. We study the *average treatment effect on the treated* (“ATT”) for units on two treatment paths, which we define as follows: consider for treatment paths \vec{k} and \vec{k}' the average treatment effect in time t on units treated with history \vec{k} as:

$$ATT(t, \vec{k}, \vec{k}') = \mathbb{E}[Y_{it}(\vec{k}) - Y_{it}(\vec{k}') | \vec{K}_i = \vec{k}]$$

This is the average treatment effect from switching to path \vec{k}' for units on path \vec{k} at time t . As discussed above, the identification problem is that $Y_{it}(\vec{k}')$ is unobservable for units with observed treatment history $\vec{K}_i = \vec{k}$.

Consider two distinct treatment paths \vec{k} and \vec{k}' , and let the first time period the treatment histories do not agree be denoted $T_{div} = \arg\min_t (k_t \neq k'_t)$. We use information from the pre-divergence periods (i.e. time periods before T_{div}) and the following assumption to impute counterfactual outcomes.

Assumption 2. [Average Bounded Deviation] Let $T_{div} = \arg\min_t (k_t \neq k'_t)$ for distinct treatment paths \vec{k} and \vec{k}' . Then the following holds for each i , $t \geq T_{div}$, any (\vec{k}, \vec{k}') pair, and some $C \in \mathbb{R}_+$:

$$\left| \mathbb{E}[Y_{it}(\vec{k}') | \vec{K}_i = \vec{k}] - \mathbb{E}[Y_{it}(\vec{k}') | \vec{K}_i = \vec{k}'] \right| \leq C \cdot \max_{t < T_{div}} \left| \mathbb{E}[Y_{it} | \vec{K}_i = \vec{k}] - \mathbb{E}[Y_{it} | \vec{K}_i = \vec{k}'] \right|$$

where $C = C(t, \vec{k}, \vec{k}')$ can depend on the time period and both treatment paths.

This assumption says the counterfactual average outcome from treating with \vec{k}' instead of \vec{k} for units with observed treatment path $\vec{K}_i = \vec{k}$ is “near” the observed outcome for units on treatment path \vec{k}' . How near? Specifically, the counterfactual average for the \vec{k} -treated group is within C times the maximum absolute difference in means for units on the two paths before the paths diverge.¹¹⁴

This assumption has “identifying power” because it relates the unobservable counterfactual outcomes—i.e. the conditional expectation $\mathbb{E}[Y_{it}(\vec{k}') | \vec{K}_i = \vec{k}]$ —to observable averages. Under this assumption, the average treatment effect on the treated

¹¹⁴Note that $C = C(t, \vec{k}, \vec{k}')$ can depend on the time period t and treatment paths (\vec{k}, \vec{k}') . In what follows, we consider fixed C -values that are constant across all post-divergence time periods. But, in principle, further assumptions on C as a function of the treatment paths \vec{k}, \vec{k}' or the time period can be incorporated. For example, in some cases there may be good reasons (e.g. theory, past studies, etc.) to believe the bound should grow in time. That is, C should be a weakly increasing function of t .

(ATT) is partially identified. That is, we may obtain upper and lower bounds on the ATT. The bounds are given by the following observable quantities for fixed C :

$$\overline{ATT}(t, \vec{k}, \vec{k}') = \mathbb{E}[Y_{it} | \vec{K}_i = \vec{k}] - \left(\mathbb{E}[Y_{it} | \vec{K}_i = \vec{k}'] - C \cdot \text{Pre-Div}(\vec{k}, \vec{k}') \right) \quad (1)$$


$$\underline{ATT}(t, \vec{k}, \vec{k}') = \mathbb{E}[Y_{it} | \vec{K}_i = \vec{k}] - \left(\mathbb{E}[Y_{it} | \vec{K}_i = \vec{k}'] + C \cdot \text{Pre-Div}(\vec{k}, \vec{k}') \right) \quad (2)$$

where the maximum pre-divergence deviation is given by:

$$\text{Pre-Div}(\vec{k}, \vec{k}') = \max_{t < T_{div}} \left| \mathbb{E}[Y_{it} | \vec{K}_i = \vec{k}] - \mathbb{E}[Y_{it} | \vec{K}_i = \vec{k}'] \right|$$

In this notation, the upper bound is given by $\overline{ATT}(t, \vec{k}, \vec{k}')$ and the lower bound is given by $\underline{ATT}(t, \vec{k}, \vec{k}')$. The method is illustrated below, with descriptive caption.

Bounded Deviation Identification (Fixed $C=1.0$)



Appendix Figure D1. Bounded Deviations (Fixed $C = 1$)

Note: In the left panel, the two lines are $\mathbb{E}[Y(\vec{k})|\vec{K}_i = \vec{k}]$ for two treatment paths: \vec{k}_1 (blue) and \vec{k}_2 (red). The black dotted vertical line is the path divergence time period (T_{div}). The blue-striped region is the counterfactual region for blue units (i.e. on treatment path \vec{k}_1) if they had instead experienced the red unit treatment path (\vec{k}_2). The maximum pre-divergence deviation is 1.69 and is shown by the black vertical line around 1925. In the right panel, y -axis is $ATT(t, \vec{k}_1, \vec{k}_2) = \mathbb{E}[Y_{it}(\vec{k}_1) - Y_{it}(\vec{k}_2)|\vec{K}_i = \vec{k}_1]$. The black dotted horizontal line is plotted for $ATT = 0$. The upper and lower bounds on the ATT for each post-1950 time are plotted as dashed-yellow and solid-green lines, resp.

D.3 Comparison with DID

For those familiar with the vast economics literature on difference-in-differences (DID), we also compare our Bounded Deviations method to the simpler DID standard. We consider here the simplest possible setting where there are only two groups—the treated and control—that may receive a single dose of a binary treatment in the year 1950.¹¹⁵ In the left panel of Figure D2, we show the standard DID strategy to identify causal effects in this simple setting. The solid-blue line is the observed average

¹¹⁵For example, this could be: “Did Congress ban the sale of good i in 1950?” If the answer is yes, then good i is part of the treated group. If the answer is no, then good i is part of the control group.

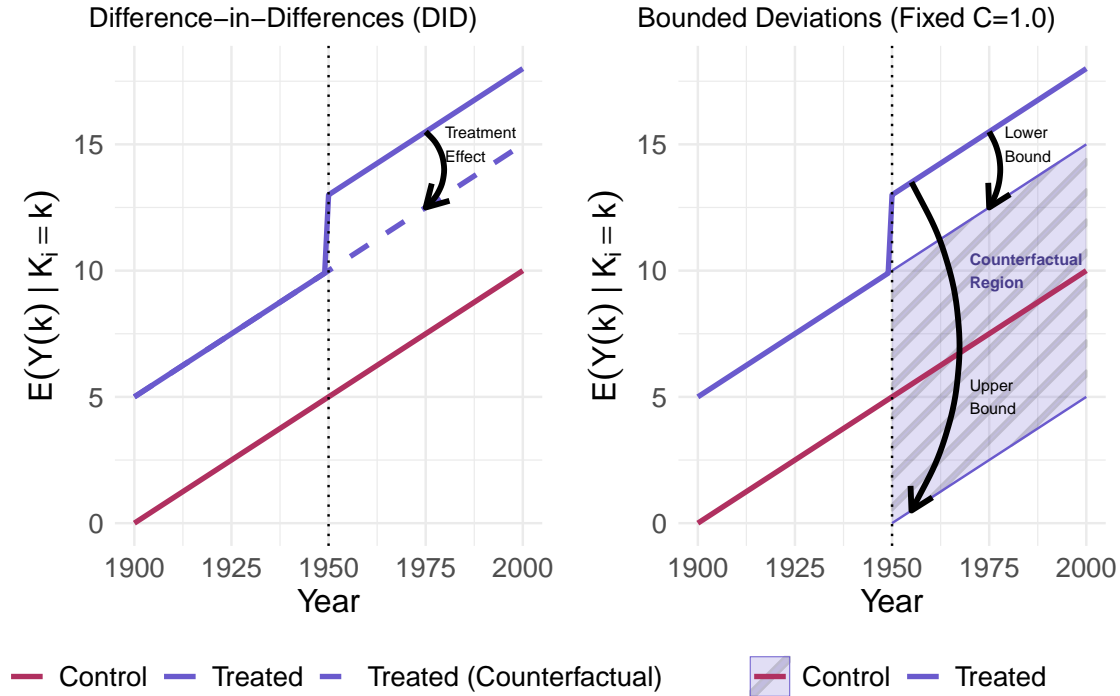
outcome in each year for treated units. The solid-red line is the observed average outcome for the control group for each year.

How to estimate the causal effect of treatment on the treated group? To estimate the treatment effect, the DID-strategy assumes that the counterfactual potential outcome would've trended in parallel to the control group. This is called the *parallel trends* assumption and it is the basis for a vast empirical economics literature.¹¹⁶ In [Figure D2](#), the dashed-blue line shows this counterfactual, unobserved potential outcome for units that actually received treatment in the left panel. This is the outcome that treated units would have experienced (on average) if they had not been treated. Notice that the dashed-blue and solid-red lines are parallel.

The treatment effect on the treated is the difference between the solid-blue and dashed-blue lines. The difference is the difference between the average treated outcome and the average control outcome for units that actually received treatment. In the left panel of [Figure D2](#), the treatment effect—called the *average treatment effect on the treated* (“ATT”)—is shown by the black arrow labeled “Treatment Effect”. Notice that in this example, the ATT is constant across the post-1950 years, but this is not a requirement of the DID strategy.

¹¹⁶See [Goldsmith-Pinkham \(2024\)](#).

Comparison of Identification Strategies



Appendix Figure D2. Bounded Deviations & DID Comparison

Note: The black dotted vertical line is the path divergence time period ($T_{div} = 1950$). The blue-striped region is the counterfactual region.

Our method in this simple setting is depicted in the right panel of [Figure D2](#). The observed average outcomes are, as before, plotted in the solid-blue and solid-red lines. The difference here is that we allow for the counterfactual (i.e. untreated) average outcome for the treated group—which before was the dashed-blue line—to not evolve in parallel with the control group average outcome. Indeed, we allow for the counterfactual average outcome to differ from the parallel counterfactual in the DID provided it remains “close” to the control group outcome.¹¹⁷ In the right panel of [Figure D2](#), this is shown by the blue-striped “Counterfactual Region”. Our method allows the counterfactual (i.e. untreated) average outcome for the treated group to lie above or below the control group average outcome (solid-red line) so long as it remains within the blue-striped region.

The treatment effect on the treated (ATT) under our method is bounded above and below. It is bounded below by taking the difference between the observed average

¹¹⁷For much more discussion of how to relax the parallel trends assumption, see, e.g., [Rambachan and Roth \(2023\)](#).

treated outcome and the largest possible counterfactual average outcome. In the right panel of Figure D2, this is shown by the arrow labeled “Lower Bound”. The upper bound is found similarly: take the difference between the observed average treated outcome and the smallest possible counterfactual average outcome (labeled “Upper Bound” in Figure D2). Finally, notice that the bounds in our example nest or contain the DID estimated treatment effect. That is, if parallel trends does hold, then the DID-based estimate is equal to the Bounded Deviations lower bound (for $C = 1$).

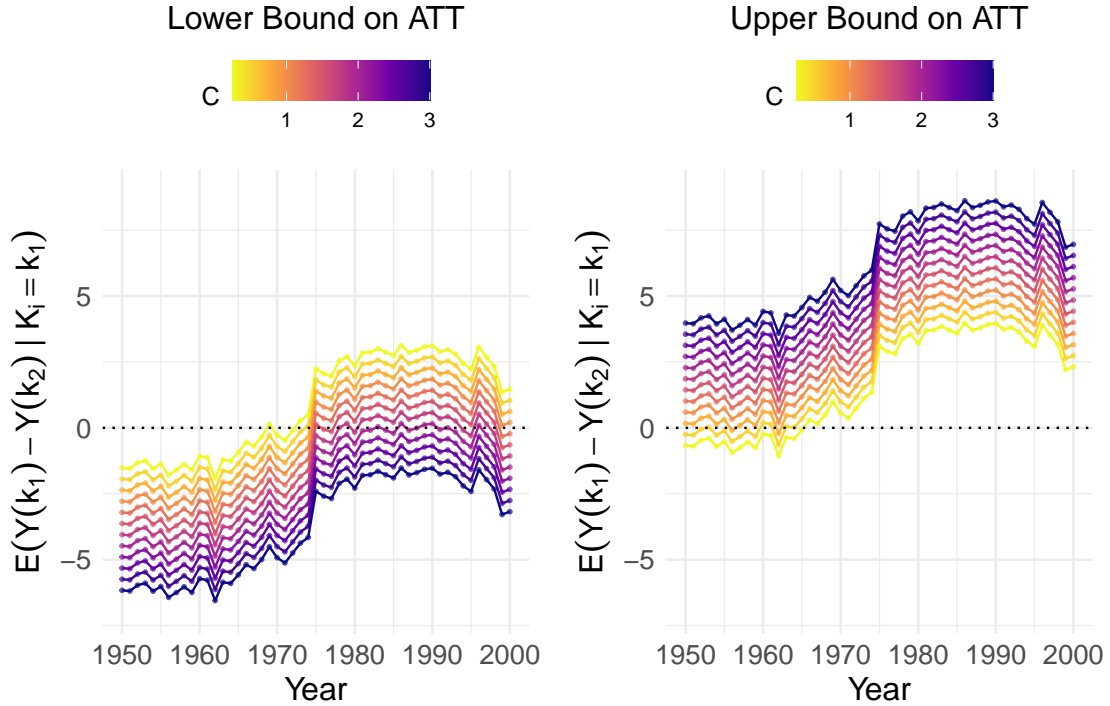
D.4 Estimation & Inference

We also compute all bounds as a function of C , which facilitates checking how conclusions change as a function of C . The estimation strategy proceeds by first partitioning the units by observed treatment paths. These paths will make up the counterfactual paths of interest. Next, for any two path pair of interest (\vec{k}, \vec{k}') we compute the divergence period T_{div} and the maximum deviation in time periods before T_{div} . We then estimate the counterfactual path using the identification assumption for different C values. Although the primary focus of the methodology here is on identification, we note that, where possible, one can estimate the conditional expectations non-parametrically and use bootstrap techniques to compute standard errors. Inference is a difficult problem, however, in cases where there is only one treated unit (as in synthetic control), so we leave further estimation and inference issues for future work.

E Appendix E: Bounded Deviations Tuning Parameter

Appendix E show how the Bounded Deviations identification strategy depends on the tuning parameter C . As shown in Manski and Pepper (2018), researchers can study how sensitive their causal estimates are to tuning parameters. We show how this identification tuning parameter affects the upper and lower bounds in Figure E1. The upper and lower bounds are plotted in the right and left panels, resp., for different values of C . In the plot, we consider varying C from being very small (yellow colored lines) to very large (dark blue lines). Smaller C values result in the upper and lower bounds being very close. For example, the yellow lower bound and yellow upper bound lines do not differ by much. By contrast, large C values result in larger counterfactual regions. This is shown by how far apart the lower and upper bounds are when $C = 3$ (i.e. dark blue lower bound and dark blue upper bound are far apart).

Bounded Deviation Identification (function of C)




Appendix Figure E1. ATT Bounds as a Function of C

Note: In each panel, the black dotted horizontal line is plotted for $ATT = 0$.

F Appendix F: Additional Results Plots

Appendix F contains additional results plots. The figures are ATT bounds plots as above for different values of the tuning parameter C . Please note in each plot the cohorts that are being compared. Unlike the ATT bounds discussed in-text, these are comparisons between treated cohorts.


Realists: Early vs Mid Cohorts



Appendix Figure F1. Realists: Early vs Mid Cohorts

Note: In each panel, the black dotted horizontal line is plotted for $ATT = 0$, the ATT upper bound is the dashed-yellow line, and the ATT lower bound is the solid-green line. The years for which the upper and lower bounds are both positive are shaded green. The years for which the upper and lower bounds are both negative are shaded red.


Realists: Mid vs Late Cohorts



Appendix Figure F2. Realists: Mid vs Late Cohorts

Note: In each panel, the black dotted horizontal line is plotted for $ATT = 0$, the ATT upper bound is the dashed-yellow line, and the ATT lower bound is the solid-green line. The years for which the upper and lower bounds are both positive are shaded green. The years for which the upper and lower bounds are both negative are shaded red.


Realists: Early vs Late Cohorts



Appendix Figure F3. Realists: Early vs Late Cohorts

Note: In each panel, the black dotted horizontal line is plotted for $ATT = 0$, the ATT upper bound is the dashed-yellow line, and the ATT lower bound is the solid-green line. The years for which the upper and lower bounds are both positive are shaded green. The years for which the upper and lower bounds are both negative are shaded red.


Originalists: Mid vs Late Cohorts



Appendix Figure F4. Originalists: Mid vs Late Cohorts

Note: In each panel, the black dotted horizontal line is plotted for $ATT = 0$, the ATT upper bound is the dashed-yellow line, and the ATT lower bound is the solid-green line. The years for which the upper and lower bounds are both positive are shaded green. The years for which the upper and lower bounds are both negative are shaded red.

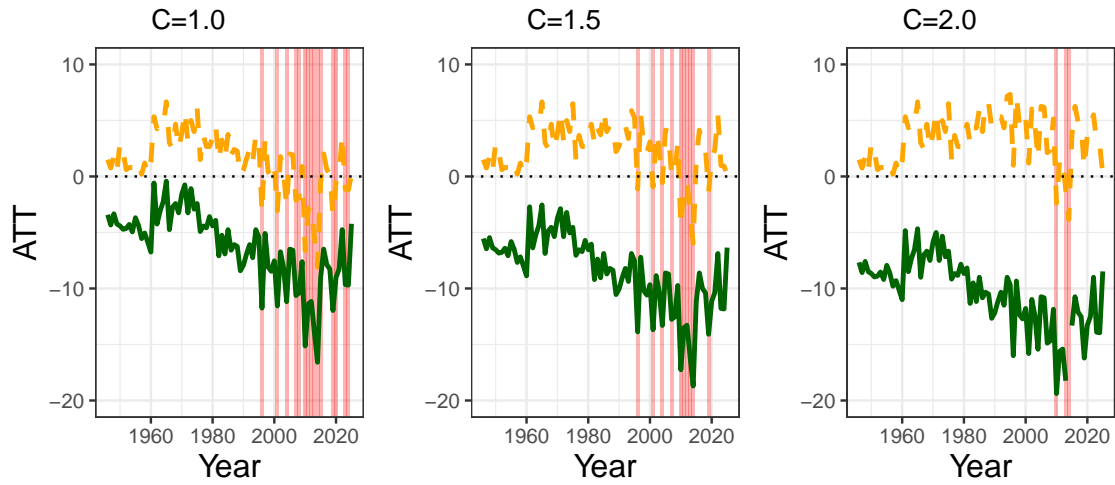
Realists: Mid vs Never-Treated Cohorts (Corbin omitted)



Appendix Figure F5. Realists: Mid vs Never-Treated Cohorts

Note: In each panel, the black dotted horizontal line is plotted for $ATT = 0$, the ATT upper bound is the dashed-yellow line, and the ATT lower bound is the solid-green line. The years for which the upper and lower bounds are both positive are shaded green. The years for which the upper and lower bounds are both negative are shaded red.

Realists: Mid vs Late Cohorts (Corbin Omitted)



Appendix Figure F6. Realists: Mid vs Late Cohorts

Note: In each panel, the black dotted horizontal line is plotted for $ATT = 0$, the ATT upper bound is the dashed-yellow line, and the ATT lower bound is the solid-green line. The years for which the upper and lower bounds are both positive are shaded green. The years for which the upper and lower bounds are both negative are shaded red.